

Project Summary/Abstract

Project Title: Coupling Sensor Networks and HPC Facilities with Advanced Wireless Networks for Near-Real-Time Simulation of Digital Agriculture

Shantenu Jha, Brookhaven National Laboratory (Principal Investigator)

Rich Wolski, University of California, Santa Barbara (Co-PI)

Douglas Thain, University of Notre Dame (Co-PI)

Mehmet Can Vuran, University of Nebraska-Lincoln (Co-PI)

Combining large-scale computing capabilities with scientific instruments and sensors (of all scales) to function as a single system has become an essential requirement for new science. Until recently, the physical and logical distance between the sensor networks and high-performance computing (HPC) facilities prevented effective connections between them due to a lack of low-latency, high-throughput, predictable communication. **We hypothesize that 5G/6G networks offer novel capabilities that can be exploited and extended to provide the low latency, high throughput, and reliability needed to perform the near-real-time coupling of edge sensor networks with simulations running in HPC facilities.**

To evaluate the proposed hypothesis, we will research and deliver software systems that will: 1) enable seamless processing of data from sensors in an Internet of Things (IoT) context with DOE leadership-class computers, 2) leverage and extend emerging 5G/6G network technologies via novel sensor slicing technology to virtualize data from both physical sensor network(s) and existing data sets, and 3) effectively compose data from the edge with data on HPC platforms. Our driving applications are from digital agriculture, a DOE priority area, for example, coupling real-time simulation of pesticide application using distributed environmental sensors communicating over a 5G/6G network to seed and update a high-fidelity atmospheric simulations on HPC.

Currently, no software effort or programming framework considers the multiscale nature of computing and devices (small edge sensors to large computing platforms). No single scale of a computing element, storage device, or sensor is sufficient. In addition, the requirements for unifying multiscale resources mandate that the system software adapts to changing conditions, failures, etc. The forthcoming advanced wireless technologies comprising 5G/6G infrastructure must include sufficient computing and storage capabilities, e.g., for low-latency decision-making, actuation, and control at the edge, particularly in remote areas that cannot be monitored or served by near-field radios. Current 5G/6G designs include “network slicing” features that enable per-flow Service Level Agreements (SLAs). However, these network-slicing technologies do not allow virtualizing end users (i.e., IoT devices). Furthermore, no 5G/6G slicing technologies or approaches are specifically architected to enable HPC coupled with real-time or near-real-time data acquisition. Lastly, large-scale batch systems, highly tuned for throughput, are not engineered to remain responsive and available in the face of fluctuating remote device and sensing availability.

Advanced wireless networks will make near-real-time synchronization of sensors and simulations possible. Even so, unexpected network events may still occur, and the system design must accommodate asynchronous operation and resynchronization as conditions dictate. We will advance and enable “end-to-end” services to create a new coupled computing environment for the computing continuum.

Our overarching research objective is to develop the software and middleware that couple computing, storage, sensing, and actuation at multiple resource scales. Such system research will be effective only if it unifies across resources and throughout the software stack. Thus, the second research objective is to investigate the software system comprehensively and end-to-end rather than as an integration of separate technologies, each developed for a different purpose. A virtualized sensor-slicing solution will be developed so the same physical sensor network(s) can provide different sensing processes without needing separate deployments. The sensor slicing concept extends the emerging 5G network slicing solutions to the individual sensing elements at the network level. Accordingly, a physical sensor can be virtualized to provide multiple

roles in each feedback loop. This innovation will seamlessly integrate *in situ* collected and locally stored data and effectively virtualize data from physical sensor network(s) and existing data sets.

Developing the end-to-end software infrastructure necessary to couple large-scale computing capabilities at one end with field-deployed sensing and actuation components at the other (while leveraging intermediary resources in between) will enable new “on-demand” applications that put HPC in the loop for critical decision-making. This multiscale, coupled infrastructure is essential for developing new applications that meet nationally critical priorities, such as biosurveillance of the national biofuel crop.

Project Title: Coupling Sensor Networks and HPC Facilities with Advanced Wireless Networks for Near-Real-Time Simulation of Digital Agriculture

Applicant Institute: Brookhaven National Laboratory, Upton, NY

Street Address/City/State/ZIP: Building 725, Upton, NY 11973

Lead PI: Shantenu Jha, Chair, Computation and Data-Driven Discovery (C3D)
Brookhaven National Laboratory
(631) 344-4780, shatenu@bnl.gov

Administrative Contact: Nicole Medaglia, 631-344-6245 medaglia@bnl.gov

FOA Number: DOE-FOA-0003300

DOE/SC Program Office: Office of Advanced Scientific Computing Research

DOE/SC Program Office Technical Contact: Robinson Pino Robinson.Pino@science.doe.gov

PAMS Preproposal tracking number: PRE-0000038808

Research area: Topic D) Advanced Wireless

Title Page Supplement

Institutions: Brookhaven National Laboratory (PI: Shantenu Jha) University of California, Santa Barbara (Co-PI: Rich Wolski), University of Notre Dame (Co-PI: Douglas Thain) University of Nebraska-Lincoln (Co-PI: Mehmet Can Vuran)

	Name	Institution	Year 1 Budget	Year 2 Budget	Total Budget
Lead PI	Shantenu Jha	Brookhaven	\$167K	\$167K	\$0.34M
Co-PI	Rich Wolski	UCSB	\$150K	\$150K	\$0.3M
Co-PI	Douglas Thain	Notre Dame	\$150K	\$150K	\$0.3M
Co-PI	Mehmet Can Vuran	UNL	\$83K	\$83K	\$0.166M

Leadership Structure: The PI and the Co-PIs each lead a research laboratory and their respective institutions that is staffed by student researchers (graduate and undergraduate), postdoctoral researchers, visiting scholars, and research developers. Project staff will be members of these laboratories but they will collaborate (under the direction of the PI and Co-PIs) to accomplish project tasks. Regular project meetings will be supervised by the PI and Co-Is who will also make all resourcing decisions.

Facilities at Brookhaven National Laboratory: The project will have access to Brookhaven National Laboratories Institutional Clusters (IC) and other high-performance and high-throughput clusters. The ICs include Annie (216 compute nodes of two CPUs Intel Xeon), Francis (a 142 compute node of Intel Xeon) and Skylake (consists of 64 compute nodes skylake processors). In addition, the project will have access to the SDCC data archive, and the Advanced Computing Lab (ACL), comprised of testbeds and platforms for research and development (R&D) of architectures and systems, as well as for siting and enabling advanced computing testbeds. Brookhaven computing facilities is connected to ESNET via dual 100 GbE networks.

Facilities at University of California, Santa Barbara: The project will have access to the RACELab at UCSB which is housed in the Institute for Energy Efficiency. In addition, the RACELab maintains a field research stations located at the Sedgwick Research Preserve in Santa Ynez, California, and the Lindcove Research Extension Center, in Exeter, California for the purpose of testing and evaluating its research results. These field research facilities will also be available to the project. The project will also have access to the Aristotle Private Cloud maintained at UCSB and the RACELab's ARM-based NVIDIA A100 GPU system and edge-edge clouds (approximately 10).

Facilities at University of Notre Dame: The project will have access to the Notre Dame Center for Research Computing (CRC) a shared campus-level facility providing approximately 20,000 cores and 200 GPUs in a variety of hardware configurations, accessible via shared batch system. The CRC also provides high performance (Panasas, Ceph) and high capacity (AFS, HDFS) storage coupled to the cluster. In addition, PI Thain operates the Cooperative Computing Lab, which provides a 10K-core opportunistic HTCondor cluster, and the DISC-II storage cluster, a 250TB SSD facility for development of experimental storage systems in support of HTC and HPC facilities.

Facilities at University of Nebraska-Lincoln: The project will have access to Nebraska Experimental Testbed of Things (NEXTT), a statewide 5G experimental testbed consisting of six sites in Eastern Nebraska Research Extension and Education Center (ENREEC), Mead, Nebraska and Lincoln, Nebraska. NEXTT is developed in a fully containerized architecture with a continuous integration/continuous deployment (CI/CD) process to automatically build, test, and release code changes as they are committed to a source repository. The project will also have access to the agricultural fields in ENREEC, including the ENREEC Spidercam facility, which is built on a 5 ac field with a Spidercam system including cameras (thermal I.R., multispectral, hyperspectral), LiDAR, and a spectrometer for advanced field phenotype research.

Student/Researcher Mentoring and Development: The project will serve to provide research opportunities for students at all levels (undergraduate and graduate) as well as visiting scholars who participate in laboratory research during their collaboration time. The project leadership team will ensure that project research fosters publication, thesis, and dissertation authorship opportunities for the students and early-stage researchers who participate. These opportunities will also include public presentations at conferences or professional meetings, poster sessions, seminar talks, etc. that will be prepared in collaboration of the project leadership team and senior personnel.

1 Introduction

Increasingly, the ability to couple large-scale computing capabilities with scientific instruments and sensors (of all scales) so that they can function together as a single system has emerged as a key requirement for new science and for new scientific impact. Mitigating the effects of climate change on agriculture and ensuring U.S. energy independence both require the modeling of, and responses to, dynamically changing physical phenomena (e.g. pathogen virility, propagation and their dependence on external conditions, etc.) that are difficult to predict. While predictions of the relevant phenomena will improve, key to this improvement is the ability to study such phenomena at ever smaller time scales in as close to real time as possible.

To bring about the necessary technological advances that will enable this next generation of science requires new computer systems research that must advance computational capabilities along two dimensions. The future technologies must be

1. *Multiscale, Multimodal, and Resilient*: they must be able to couple and decouple computing and storage resources with instruments, sensors and actuators at all scales. No one scale of computing element, storage device, or sensor is sufficient. Further, the necessary computational and storage capabilities must be designed for long-lived, unattended operation in the face of changing conditions. This resilience must be a fundamental design requirement that unifies all levels of the system software stack at all resource scales.

2. *Adaptive*: The requirements for both unifying multiscale and multimodal resources mandate that the system software adapt to changing conditions, capabilities, usage modalities/QoS constraints, failures. This necessitates the importance of agile and adaptive resource management and workload execution properties.

In this work, *we hypothesize that 5G/6G networks offer novel capabilities that can be exploited and extended to provide the low latency, high throughput, and reliability needed to perform the near-real-time coupling of edge sensor networks with simulations running in HPC facilities.* More specifically, the coupling of large-scale systems with scientific instruments, sensors, and actuators at all scales requires a new approach to adaptive workflow management and new system software abstractions. In particular, this coupling requires that components describe and expose performance and resilience properties in a way that new scheduling and software infrastructure can and must exploit though novel scheduling algorithms and distributed system software.

We propose to investigate this hypothesis in the form of XGFABRIC – an end-to-end, multi-scale, and adaptive distributed system that couples leadership class resources with scientific instruments, sensors, and actuators through a tiered network of supporting computational and storage resources. Central to XGFABRIC are new workload partitioning and scheduling approaches that enable scale, adaptivity, and resilience for applications using coupled and multiscale infrastructure.

Our research approach will be empirical and artifacts-based. We plan to develop XGFABRIC as a evolving prototype that permits continuous evaluation of the research results this project will generate. We will also use “real-world” applications focused on digital agriculture as driving applications. We will deploy XGFABRIC in the field using sensing, edge, and high-performance computing resources as part of our efforts to evaluate its effectiveness.

Multiscale, Multimodal Systems

The state-of-the-art computational science focused on digital agriculture, its dependence on climate change and its relationship to energy independence (e.g. carbon neutrality, renewable energy sources, etc.) requires the most powerful leadership class machines that are available. However, because the phenomena that are critical to understand and predict are changing more dynamically, the resolution with which these phenomena must be measured and modeled is becoming ever finer.

While the need for high-resolution computational science is rightfully associated with the twin needs for data and computational scale, when it is driven by increasing dynamics, scale alone is insufficient. For example, in a disease propagation response context, the speed with which the data can be acquired, the resilience of the infrastructure (hardware and software) with respect to the pathogen itself, and the ability to

adapt computation and data management to change conditions emerge are essential properties that must be part of the overall system that is responsible for delivering data and computational scale.

At present, no software effort takes these requirements to be essential in this computational context. The cloud, for example, provides no resilience guarantees (or even service level objectives) for edge devices or sensors in an “Internet-of-Things (IoT)” context. The coming cellular technologies that comprise 5G and 6G infrastructure do not include sufficient computing and storage capabilities, particularly in remote areas that must be monitored but cannot be served by near-field radios. Large-scale batch systems, highly tuned for throughput, are not engineered to remain responsive and available “on-line” in the face of changing phenomenological dynamics and fluctuating remote device and sensing availability. Our work will advance and enable resilience “end-to-end” to create a new coupled computing environment for multi-scale, dynamic, and “streaming” science applications.

The coming advances in 5G and 6G connectivity will make it possible to build and maintain multi-scale, networked computational infrastructure that is capable of interfacing the smallest of devices at the edge with a spectrum of larger resources up through, and including, the cloud and leadership-class HPC systems. These computational conglomerates are *multi-scale* because their constituent components span several orders of magnitude in computational and storage capability. Devices at the extreme edge of the network that interface with low-power, highly durable sensors and actuators may only support 10s of kilobytes of memory with single-core CPUs clocked at 10s of megahertz. These devices must interface to successively larger small board computers (e.g. low-power Linux machines), edge-clouds [8], regional and private clouds, and (at the largest end of the scale), public clouds and HPC resources. These systems are also *multi-modal* in that the function of all of the resources (including device and device controllers) must be able to change in response to the needs of the application components they host. For example, a sensor controller interfacing with a physical sensor may need to also host (temporarily or permanently) telemetry “filtering” software to reduce the network message payload size and thereby conserve battery power.

The need to provision computational resources at multiple scales and to manage the modes in which these resources operate (particularly at the edge) gives rise to a new model for distributed systems that couples physical sensing and actuation with amalgamations of 5G and 6G interconnected resources. To leverage 5G and 6G technologies for multi-scale, multi-modal systems requires a *digital-physical fabric* – a computational infrastructure in which the resources comprised by the system are *interwoven* by a common and unifying software infrastructure. Digital-physical fabrics are distinct from “digital twins” [32] and cyberphysical systems [4] in that they do not separate and categorize the resources as either digital or physical. These earlier paradigms focus on managing the interface between physical “real-world” interactions and computational capabilities – the digital and the physical are separate but equal. Alternatively, a digital-physical fabric blurs these distinctions through its software infrastructure so that the system functions as architected, end-to-end, from the physical interactions through final data analysis. In particular, it manages both the spectrum of scales that comprise the end-to-end system and the panopli or modes that the resources may assume. Indeed, digital twins and cyberinfrastructure can be important components of a digital-physical fabric – we describe the incorporation of a digital twin for biosurveillance and biosecurity in Section 2.1.

Our work described in this proposal will explore the development of digital-physical fabrics using 5G and 6G technologies as the digital fabric that makes a unified and durable end-to-end system possible.




Figure 1: In the digital continuum, observational and sensing data flows from the edge to the center. Processed data and updates to AI/ML models flow from the leadership class platform to the edge and even to the field. (Adapted from DOE 5G Enabled Energy Innovation Workshop Report [5])

2 Science Drivers for Advanced Wireless for Digital Agriculture

2.1 SD1: Citrus Under Protective Screens

The citrus production industry is currently developing remediation strategies for the Asian citrus psyllid which carries the huanglongbing (HLB) “citrus greening” disease. HLB has devastated the commercial citrus industry in Florida and Texas with an annual cost of more than \$1B US [24]. From a biosafety perspective, HLB is a significant vector. Pesticides and disease-resistant cultivars have, so far, proved ineffective. Its effect on citrus production in the south has been rapid and irreversible.

In California, where the disease is present but not yet epidemic, growers are experimenting with siting orchards inside large, protective screen houses. The Citrus Under Protective Screening (CUPS) project is an at-scale pilot for screen-house citrus production located at the Lindcove Research Extension Center in Exeter, California. While CUPS is specifically testing HLB control, it represents an approach that is effective against any insect-born pathogen for which typical husbandry practices are ineffective.

The goal of CUPS is to understand the growing environment and commercial agricultural viability of screen-house citrus production. Citrus trees have useful production lifetimes that exceed 20 years. CUPS is effective as long as the trees that are introduced into the screen house are disease free and the screen remains in tact. For commercial viability, the screen houses must be large (covering several acres each) and they must accommodate tree canopy and harvesting equipment that require 25 to 30 feet of vertical space.

Detecting and rapidly repairing screen breeches in commercial scale CUPS is a critical open problem. While industrial accidents that cause screen damage can be detected and rapidly reported by workers, unobserved events (e.g. bird strike, foraging fauna, damage concomitant with theft, etc.) can cause screen breeches that must be detected.

Our team has been working to instrument and analyze the growing environment within the at-scale CUPS structure in Exeter. As part of that on-going work, we have developed a Computational Fluid Dynamics (CFD) model that can model to predict airflow within a CUPS screen house in near real time based on instantaneous wind, temperature, and humidity measurements taken and the screen boundaries (both inside and outside). Analytically, the goal of the model is to provide growers with decision support for input events such as pesticide or fertilizer spraying, frost prevention, etc. where the grower must make a decision regarding timing, location, and quantity of input to apply.

However, we are also exploring whether the model can detect screen breach. Specifically, once the model is calibrated, a deviation between predicted and measured airflow can portend a possible screen breach and, perhaps, an area of the structure where the breach may have occurred.

Note that we plan to structure the coupling of real-time sensor data with CFD as a “digital twin” in which the true atmospheric conditions within the structure are “twinned” by the results of the CFD model for the interior of the structure. The model results will inform both modality changes in the sensing infrastructure and data calibrations (back tested against historical data) that are necessary to maintain model accuracy.

Our team will also be deploying a Farm-NG [13] wheeled robot with autonomous-driving capability within the CUPS structure. As a driver for XGFABRIC research, our plan is to investigate whether it will be possible to detect a potential breach (using a large-scale HPC machine to run the CFD model which is parameterized by real-time *in situ* boundary conditions), compare the modeled airflow to measurements taken for the same time period within the structure, and if they do not match, dispatch the robot to surveil the region of the screen where a breach may have occurred using an on-board camera. The XGFABRIC digital-physical fabric will incorporate robot-based sensing and robot route planning, thereby linking it to, and augmenting the CFD-based digital twin for the screen structure.

This ambitious application illustrates how a digital-physical fabric can enable new biosecurity capabilities. However, to bring it to fruition requires the ability to amalgamate computational resources at all scales, and to “close the loop” between sensing, computing and storage, and actuation.




Figure 2: The National Virtual Biosecurity for Bioenergy Crops Center (led by Brookhaven National Laboratory) is tasked with designing the integrated infrastructure to support the collection of field data, detection of pathogens, and sampling of bioenergy crops with traditional simulation-modeling capabilities.

2.2 SD2: Biosurveillance and Biosecurity of Bioenergy Crops

The US and other nations have set net-zero carbon goals, which require the development of comprehensive and scalable technologies to replace fossil-carbon energy sources and remove atmospheric carbon through net-negative emission technologies. Expanding the US Bioeconomy is part of the US Government’s “all of the above” approach to meet its net-zero goals (US Long-term Climate Strategy, 2021). The US aims to build the US bioeconomy, which includes enhancing biosafety and biosecurity to reduce risks to the US bioeconomy. In an expanded bioeconomy, a significant fraction of the agricultural sector will be devoted to growing bioenergy crops to support the production of biofuels and bioproducts. The emergence and spread of diseases among bioenergy crops could significantly impact yields and destabilize the US bioeconomy. However, with the projected expansion of bioenergy crops, the risks to crop health associated with the emergence of new diseases, the climate-driven spread of known diseases into new regions, and the accidental or deliberate release of plant diseases need to be considered. In this context, based on experimental and computational facilities, new capabilities are needed to detect, characterize, model, and mitigate biothreats to bioenergy crops.

There are unique challenges in handling extensive remote agricultural field data in real-time streaming analysis with a performance guarantee. Notably, it requires: 1) reliable network connectivity to transmit large volumes of field phenotype data (i.e., drone flying with hyperspectral imaging), 2) intelligence surveillance planning (i.e., optimal scanning strategies), and 3) efficient real-time streaming processing at the edge. (2) and (3) often must be coupled to traditional modeling and simulation capabilities.

These workflow challenges pose an opportunity to incorporate and improve 5G or upcoming 6G wireless communication to enable (1) high-throughput transmission with low latency and large coverage area, (2) autonomous data curation/collection using AI or OED (Optimal Experimental Design), and (3) real-time edge processing capabilities.

Advanced software services and analytics must address additional challenges to play an essential role in the biosecurity of bioenergy crops. These include model validation, sparse datasets that cause ML training difficulties, and integrating data from different modalities to create interoperable datasets. Data processing capabilities at the edge must be aligned with traditional simulation and modeling capabilities to create a unified, integrated modeling approach for bioenergy crop security.

2.3 SD3: Real-time Soil Monitoring and Irrigation Automation

Eighty percent of the 24 million Americans without high-speed Internet live in rural areas [3] and farm to feed the world's population. The world population is expected to reach approximately 9.1 billion people by the year 2050 [34], which requires a 70% increase in food production with minimal adverse environmental impacts. Rural broadband connectivity, both on-farm and off-farm, is key to a digital agricultural infrastructure to improve U.S. food security [17].

Despite connectivity challenges, soil moisture measurements have informed irrigation decisions for decades. However, real-time measurement and automated irrigation solutions still face challenges. The challenges include a lack of robust rural wireless coverage and difficulty installing and removing sensors before and after a growing season. The need for real-time in-situ information from large and diverse agricultural fields has given rise to **Agricultural Internet of Things (Ag-IoT)**, which could provide in-situ monitoring capabilities (e.g., soil moisture, temperature, electrical conductivity) and, when interconnected with existing field machinery (seeders, irrigation systems, sprayers, combines), enable field autonomy. Field autonomy through autonomous irrigation scheduling, variable rate irrigation to water different parts of the field with different amounts of water, and autonomous fertigation, will pave the way for better food production solutions [39]. One of the biggest challenges facing precision agriculture is getting information to growers and enabling them to take action. To this end, there is a growing need for in-field sensors, aerial or satellite imagery, communication devices, cloud-based analytics, remote monitoring, and real-time control of in-field equipment. A high-bandwidth, wide-area network with good penetration of foliage, buildings, and terrain is needed to address the challenges of limited or nonexistent coverage in rural areas.

Timely and high-resolution insights into soil characteristics are challenging due to the difficulty of seasonal installation and removal of soil sensor systems, along with the lack of cost-effective and high-performing data acquisition. Coupled with the recent, reliable, high bandwidth wireless connectivity in rural communities, these systems will contribute significantly to the adoption of technology in production.

The interplay between digital agriculture and rural broadband connectivity is best exemplified in existing scientific and commercial practices. For example, at academic labs (e.g., UNL's East Nebraska Research Education and Extension Center (ENREEC) in Mead, Nebraska), scientists and researchers utilize advanced sensing equipment (e.g., hyperspectral sensors mounted on a spidercam structure or unmanned aerial vehicle (UAV)-mounted high-resolution cameras) to collect large volumes of data (Fig. 3). However, this data is only stored in an SD card and manually transferred to a data center for post-processing. The lack of rural broadband connectivity prevents real-time access to these data and stifles scientific understanding and progress. There is a need to transform emerging 5G technology verticals (e.g., network slicing) for agricultural requirements such that unique characteristics of agricultural fields (e.g., limited wireless infrastructure and coverage, unique built infrastructure, unique wireless propagation characteristics due to dynamic changes in crop canopy height and field operations) are leveraged.



Figure 3: UNL Spidercam field phenotyping facility.

3 Work Track 1: Digital-Physical Systems

The innovations in terms of connectivity that 5G and 6G bring militate for a new approach to closing the loop between sensing, computation and storage, and actuation. We term the unified hardware and software system that allows computational and storage resources at all scales to act in a coordinated way a *digital-physical* system. Ubiquitous, high-quality, and provisionable connectivity makes it possible to consider digital and physical interactions within the same system architecture.

With Track 1, we will develop a multi-scale, distributed digital-physical fabric to unify components in an end-to-end deployment. We will investigate a “full-stack” approach, which consists of language-level, runtime-system middleware, operating systems-level, and networking innovations, to unify a heterogeneous, distributed, and multi-scale set of resources as a tiered end-to-end system. Note that our approach will combine fully native “bare metal” deployments, with middleware that can function as part of an existing software ecosystem. SensorSlicer (Section 3.1) will provide QoS-enabled provisioning capabilities for end-devices that the “FabriStack” (Section 3.2) will unify to create an end-to-end application platform.

3.1 SensorSlicer: End-to-end Network/IoT Slicing for Virtual Data Collection – Lead: UNL

We will leverage and extend emerging 5G/6G network technologies via novel sensor slicing technology to virtualize data from both physical sensor network(s) and existing data sets.

Research Challenge: The forthcoming advanced wireless technologies comprising 5G/6G infrastructure must include sufficient computing and storage capabilities, e.g., for low-latency decision-making, actuation, and control at the edge, particularly in remote areas that cannot be monitored or served by cloud-based core sites. The envisioned suite of science

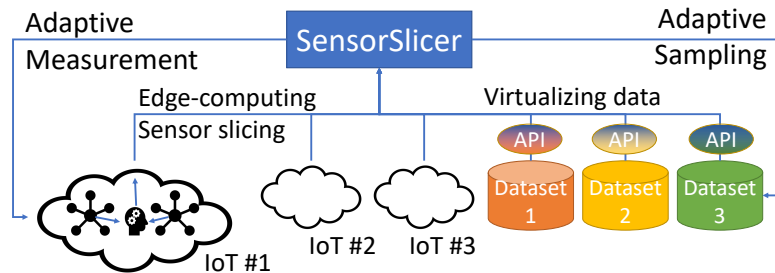


Figure 4: SensorSlicer: Virtualizing sensors in IoT networks and datasets for cost-effective data collection and adaptive measurement.

drivers (Section 2) emphasizes the crucial need for robust, reliable connectivity solutions in farms. More specifically, swarm robotics for efficient citrus farm management necessitates seamless, reliable, low-latency connectivity for real-time coordination and control. Similarly, hyperspectral images from drones and real-time surveillance data flow require high-capacity data transfer and low-latency connectivity for effective coupling with modeling and simulation tools. Moreover, emerging Ag-IoT networks with numerous end-points require field-level coverage with high-density and energy-efficient connectivity. Integrating such diverse use cases with specific per-flow connectivity requirements, highlights *the central role that advanced network infrastructure plays* in realizing the potential of HPC with scientific instruments and sensors.

Furthermore, predominantly in rural areas, farm fields suffer from limited infrastructure. Integrating diverse use cases in a single farm is cost-prohibitive with existing approaches. More specifically, current solutions require either multiple sensor networks to be deployed for multiple use cases, increasing the cost per use case; or sharing the data collected by a sensor network at the cloud end, which risks privacy. To this end, virtualizing data from both physical sensor network(s) and existing datasets is a major challenge.

Network Research: Current 5G/6G designs include *network slicing* features that enable per-flow resource allocation such that networking, computation, and storage elements can be assigned to each type of traffic flow in a more granular fashion. Thus, we virtualize a single physical network infrastructure using multiple network entities tailored to each type of traffic flow (e.g., high throughput, low latency, high density). Network slices are structured based on Service Level Agreements (SLAs), according to which resources are allocated to an end-point to meet SLA requirements. However, emerging 5G network-slicing technologies do not allow virtualizing end users (i.e., IoT devices) because they are primarily geared towards general-purpose networks, where the network infrastructure provider do not have access to end-point devices (e.g., cellphones). However, further virtualization of the IoT end devices is possible in more specific, purpose-driven networks, such as agricultural IoT. Accordingly, *the same physical sensor device could be virtualized to act as different sensors for different use cases*, significantly decreasing deployment costs.

SensorSlicer abstraction along with 5G network slicing technology (Fig. 4) provide *several advantages for seamless integration of in-situ measurements, established datasets, and scientific computing*: (1) The

same physical sensor network deployment could be utilized for different workflows, (2) Workflow tasks could be virtualized deployed anywhere throughout the field-edge-network-cloud path, and (3) Adaptive data collection methods, which fuse in-situ sensor sampling and database sampling, could be developed. To this end, we will focus on two major research tasks.

SensorSlicer: At early stages of the project, we will deploy a 5G standard-compliant private cell using existing well-maintained open-source radio access network (RAN) and core software packages (e.g., OAI [1], srsRAN [2]). These packages allow connected common-off-the-shelf (COTS) endpoints (e.g., 5G dongles) and fully open-source software-defined radio-based endpoints while allowing innovations in the 5G stack. Accordingly, the network slicing functionalities will be tailored to the field-to-cloud workflow management using COTS endpoints. Then, a virtualized sensor-slicing solution will be developed using an open-source SDR-based endpoint implementation such that the same physical sensor network can provide different sensing processes without the need for separate deployments. The sensor slicing concept extends the emerging network slicing solutions at the network level to the individual sensing elements. Accordingly, a physical sensor can be virtualized to provide multiple roles in each workflow. In the second part of the project, we will work on providing an adaptive sampling functionality on a per-flow basis such that workflow-adaptive sampling could be supported, leading to an end-to-end feedback loop for field-to-cloud workflows as we describe next.

Field-to-HPC Adaptive Sampling: In sensor-aided workflows, current practices involve in-situ measurements that are taken with prior assumptions to build an associated model. However, validating these assumptions or improving resulting models is challenging because our understanding of the datasets evolves. Existing practices of deploying sensor networks for each workflow implies a considerable footprint in terms of energy consumed for sensing, communication, processing, and storage. Therefore, we will develop workflow-aware adaptive data measurement solutions as a key component of SensorSlicer. Specifically, we will develop novel sensing systems for adaptive sampling in agricultural fields. Relying on our decade-long experience in deploying Ag-IoT systems [26, 27], we will develop new sensors that are capable of real-time adaptation and edge-based adaptive sensing. Our edge platform (FabriStack) will deploy lightweight AI-based sampling processes in-situ such that data transfer and cloud storage can be minimized without affecting information fidelity. Since computing processes are virtualized, processes can be deployed across the field-to-HPC continuum to take advantage of these new sensing and network slicing capabilities.

Evidence that this approach will succeed: SensorSlicer relies on 5G network slicing concepts that are a part of the recent 5G standards. More importantly, recent open radio access network (Open-RAN) efforts, which aim to softwareize RAN deployments through open source software and commodity hardware, have resulted in standard-compliant implementations of 5G architectures (e.g., OAI [1], srsRAN [2]). Our group has extensive experience deploying these packages in large-scale wireless experimental testbeds [20].

Nebraska Experimental Testbed of Things (NEXTT) is a state-wide remotely accessible advanced wireless infrastructure developed in collaboration with the City of Lincoln and UNL [41]. The testbed includes 5G sites deployed on a Lincoln traffic intersection (AVC site) and four rooftop sites on the UNL campus as well as an indoor site. A sixth site was deployed and operationalized at ENREEC agricultural fields, replicating this site design. High-end sub-6GHz MIMO software-defined radio (SDR) transceivers are connected to UNL's Holland Computing Center (HCC) through 20Gbps dedicated fiber links. NEXTT has been designated an FCC program experimental license for sub-6GHz and mmWave bands.

NEXTT employs a fully containerized architecture with a continuous integration/continuous deployment (CI/CD) process to automatically build, test, and release code changes as they are committed. CI/CD and containerization simplify integration of existing open-source RAN packages in a modular fashion. To date, we have tested integration of different combinations of srsRAN, OAI's EPC, and RAN packages.

In this project, we will build on the established 5G sites [41], fully containerized software architecture for 5G RAN [20], and our decade-long experience in deploying remotely accessible agricultural sensor net-

works with autonomous irrigation systems [7, 27, 28, 33, 39]. The proposed SensorSlicer architecture relies on established network slicing solutions and requires extensions of this approach to resource-constrained IoT devices. We will utilize open-source implementations of 5G endpoint architectures to implement SensorSlicer and validate the proposed solutions in real-life agricultural field deployments. By doing so, SensorSlicer will coherently integrate with the full-stack system software (FabriStack) that we describe next.

3.2 The FabriStack: System Software for Digital-Physical Deployments – Lead: UCSB

The XGFABRIC system software must support a common set of software abstractions across all resource scales required by an application that couples devices, instruments, moderate computing and storage resources, and large-scale computing facilities. Further, these abstractions must be defined to support automated deployment, performance, and security at all scales. Finally, the system must be usable by scientific programmers and developers (i.e., the abstractions cannot be so complex that only a handful of experts can use them).

To achieve this unification, we propose a layered software architecture. Figure 5 shows the relationship between the systems components we propose (shaded in the figure), multi-scale resources (along the bottom), existing systems, technologies, and applications (along the top). This *FabriStack* consists of 3 interoperating systems that enable applications to transparently span multi-scale, heterogeneous resources: CSPOT, Laminar, and Glide In.

CSPOT investigates multi-scale (sensors-edge-cloud) systems software for IoT deployments in remote locations where there is no access to the electrical power grid or networking [8, 14]. Laminar is a dataflow program representation that uses CSPOT as a runtime system. It is an innovation we propose to ease the programming burden associated with developing CSPOT applications natively. Finally, Glide In (a term borrowed from HTCondor [12]) is a facility for dynamically deploying and then decommissioning CSPOT to and from a large-scale, batch-controlled HPC systems. Together, these advances form a set of abstractions and runtime services that span the digital continuum that can be treated by developers and operations personnel as a single unified platform. Metaphorically, the FabriStack “weaves” the continuum of digital “threads” into a single, consistent “fabric” using 5G/6G connectivity.

3.2.1 CSPOT – Serverless Platform of Things written in C (UCSB)

The Research Challenge: To allow an application to amalgamate devices, computers, and storage at all resource scales (from embedded systems to supercomputers) requires a common set of unifying runtime system and programming abstractions. The alternative of integrating separate technologies developed for disparate purposes (e.g., embedded systems and cloud-web service for e-commerce) yields poor security, poor performance, a lack of maintainability, and low programmer productivity. Our team used the commercial cloud IoT technologies extensively in agricultural settings. These applications require between 7 and 15 separate code stacks and services to work together to enable basic computation and storage. Further, porting these technologies to new devices can be laborious. For example, an AWS-compatible implementation of TLS for a new and particularly low-power microcontroller took more than 7 weeks by an expert embedded systems programmer working in our group.

Addressing this challenge is difficult because it is difficult to “miniaturize” abstractions developed for full systems (and collections of full systems, like clouds) so that they work efficiently at *all* resource scales, including the small ones. Our previous work, and work proposed in this effort addresses this challenge.

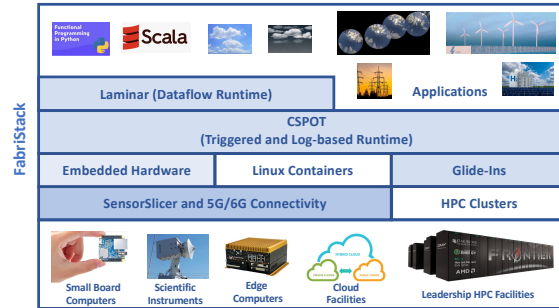


Figure 5: The Proposed FabriStack Layered Software Architecture. Shaded components comprise the FabriStack. CSPOT provides triggered computation and log-based storage universally – natively and as a guest in containerized or virtualized environments, as middleware on native operating systems or as a “glide in” on batch systems. Applications (at the top of the figure) use either Laminar dataflow programming or the CSPOT API.

CSPOT research to support XGFABRIC: For XGFABRIC, our research will focus on adapting CSPOT to use SensorSlicer to implement QoS and networking optimization in a network transparent way. Currently CSPOT implements network transparency using URNs to name its abstractions and all network performance optimizations and resilience functions are implemented by the CSPOT protocols. SensorSlicer will make it possible to create QoS-aware CSPOT logs and to tune internal resilience parameters (e.g. message timeouts). This integration not only extends the reach of CSPOT through its ability to leverage 5G and 6G connectivity, it improves its ability to scale out, through better timeout management and latency guarantees, and down, by using QoS guarantees to minimize log sizes.

This latter optimization is especially needed at the device level where memory is often severely restricted and persistent storage carries a heavy power cost. For embedded sensor controllers, the operation of the network via radio is often the largest part of the power budget allocated to each duty cycle. SensorSlicer will make it possible to offload the responsibility for implementing QoS to the network infrastructure, thereby reducing the network processing load on the device. We will integrate this capability transparently in the FabriStack using CSPOT as the integration vehicle.

Evidence that this approach will succeed: CSPOT imposes two unusual constraints on its programmers. The first is that all persistent storage must be log based. Append-only storage semantics have long been the basis for scalable, replicated storage (typically using eventual consistency replica protocols [22]) which has utility in an IoT context. However, log-based storage adds to this scalability strong crash-consistency semantics and good network partition recovery properties, both of which are essential for IoT. The second property is that computation can *only* be triggered when data is appended to some CSPOT log (called a *WooF* in CSPOT parlance). This property ensures that CSPOT (which uses a log to dispatch all computations internally) captures causal ordering of events automatically. That is, the CSPOT dispatch logs are used to record all append operations and scheduled computations making it possible to determine, as ground truth, causal execution dependencies.

These two peculiar requirements, originally introduced to support IoT in remote “infrastructureless” geographic locations, have important ramifications for XGFABRIC. First, CSPOT is designed specifically to address the tension between resilience and performance. All CSPOT functions execute concurrently and the only synchronization is the assignment of a sequence number to a log append. As a result, regardless of the degree of computational concurrency, a power failure or a crash leaves the logs in a consistent state (much like in a log-structured file system [25]). Similarly, a log append either fails to occur (in which case it will be retried) or it succeeds and the assigned sequence number in the response is lost (in which case a retry may increment the sequence number for a duplicate append). CSPOT client-side generated sequence numbers can then be used to implement duplicate suppression.

A second attractive property is that logs can be implemented very efficiently (more efficiently than file systems) at all resource scales. For example, CSPOT runs on an ESP8266 [11] microcontroller with 112k bytes of memory and an 80Mhz clock designed for embedded systems. A CSPOT Put () operation (the API command to append to a WooF) requires approximately 26 microseconds on this device [40]. This same simple approach (and API) scales up to larger computers and, ultimately, to clouds (where CSPOT is 1 to 3 orders of magnitude faster than FaaS platforms) creating a unified way to implement efficient resilient computation at *all* resource scales.

3.2.2 Laminar – Dataflow Programming for CSPOT (UCSB)

Research Challenge: CSPOT programs are notoriously difficult to write and to debug, particularly for programmers unfamiliar with development for distributed execution. While the CSPOT Put and Get operations are simple and act as synchronous functions, handler invocation is asynchronous and concurrent. Further, the only synchronization mechanism available is the transaction that assigns a sequence number (which is exposed through the Put/Get API) to a log entry when it is appended. Adding to this source of difficulty, CSPOT’s append-only semantics essentially creates a new version (indexed by sequence number)

of each state update in a program each time the variable is updated. When program variables are structured or aggregate, the tendency is to serialize them in the log. Thus, the programmer familiar with typical concurrency management (e.g. locks, semaphores, condition variables) and mutable data structures often write “naive” CSPOT programs that “spin” polling for a sequence number increment as a synchronization mechanism, and log scans (to find the latest update to an aggregate data structure).

Addressing this challenge has proved difficult for several reasons. First, much of CSPOT’s performance and all of its multi-scale portability (to the low end) accrue to its simplicity. It is analogous to an “assembly language” for building multi-scale, distributed applications end-to-end.

Programmability Research to Support xGFABRIC: We believe that a dataflow programming model can significantly enhance programmer productivity for applications that span multi-scale systems. To enable this, we have designed Laminar as a dataflow intermediate representation that both encodes dataflow computations (expressible as Directed Acyclic Graphs or DAGs) in CSPOT WooFs and also implements a dataflow runtime execution of these computations. In particular, the Laminar runtime implements efficient synchronization between data availability and handler invocation on behalf of the programmer.

We represent complex programs in Laminar as a hierarchy of DAGs where each DAG takes a set of inputs and produces a set of outputs. This hierarchical structure (inspired by IF1/IF2 [29, 30]) has two attractive properties. First, it supports program composition and modularity. A subgraph can be treated as a node within the graph in which it is embedded. Secondly, the graph hierarchy can encode functionality that is difficult to represent as a DAG. Iteration expressed a “for loop,” for example, can be represented as a compound node consisting of three subgraphs. While powerful, dataflow necessarily elides architectural details at the application level. A developer cannot reason about “where” each dataflow node will fire nor make assertions about “when” it will execute once it is enabled by the presence of its inputs.

We plan to address these restrictions in two ways. First, Laminar programs consist of a “preamble” that specifies the application’s deployment and the “body” which implements the application’s logic. The preamble is generated using the same code generator as the body, but the code that is created is sequential. It initializes CSPOT WooF’s to encode the deployment information (e.g. URN’s associated with each CSPOT log) and the connectivity information. Once distributed, the preamble need only be executed once and the resulting deployment state is available to multiple executions of the application body.

Our plan is to enhance the preamble with language pragmas that allow a developer to attach QoS and capacity requirements to each node and edge in a Laminar program. These requirements will allow the Laminar runtime to select appropriate execution sites and provision the necessary QoS from SensorSlicer for the intervening network connectivity. For example, in a case where functionality can be implemented either on resource restricted devices (say one Laminar node per device) or (with a higher latency) on an edge computer, in the event of device failure(s), it will be possible to combine the nodes from the failed devices into a “supernode” that can be executed on the edge computer without changing the semantics of the Laminar program, and to provision the necessary network capacity for the supernode using SensorSlicer.

We also plan to use this functionality to incorporate large-scale computing facilities into Laminar programs. Rather than expressing all computations at the finest level, we envision special “macro nodes” that embody whole programs (e.g. written in C, C++, or Fortran using MPI) that will execute on systems where xGFABRIC must act strictly as a guest (i.e. is operating as middleware) as on the LCF. These macro nodes act as nodes within the Laminar program that cannot be decomposed (at least, automatically) into finer-grained nodes (although other nodes may be combined with them should they too require execution on a large system). Our Laminar benchmark suite includes both Map-Reduce and Logistic Regression implemented as macro nodes with embedded parallelism. The high-level workflow is triggered and managed by Laminar and while the embedded computations within each node can either be Laminar primitives or programs written in high-level languages that gather their inputs and return their outputs through Laminar.

In this way, Laminar becomes a high-level control program that unifies execution across a multi-scale

resource deployment, where it can execute either on “native” CSPOT or as a middleware interfacing to “macro nodes” that do not use CSPOT as a runtime. Because dataflow implements applicative program semantics, we also plan integrations with popular high-level languages such as Functional Python and Scala. It should also be possible to using graphical tools (such as Graphviz [9]) with Laminar to aid in program development and debugging.

Evidence that this approach will succeed: The recent rise and popularity of dataflow programming in cloud computing (e.g. Google Cloud Dataflow and Microsoft Analytic Dataflows, Node-RED, GNURadio) show that dataflow can simplify programming and management of complex applications while facilitating reuse and customization. Our work will bring these benefits to multi-scale, distributed, and failure prone settings by combining them with a foundation of portability and resiliency at the runtime layer.

4 Work Track 2: Edge-to-Exascale Resource and Workload Management

Research Challenge: The principal workflow challenges of task placement, resource shaping and selection, must be addressed in the context of dynamic and multiscale resources, and application and resource level heterogeneity.

Multiscale resources vary in terms of availability, capacity, and response over several orders of magnitude. Together with application and resource heterogeneity, an adaptive response based on a “one pass” translation of a task-graph into an execution plan is rendered ineffective. The workflow control system must provide translation via multiple intermediate representations. The specifics of translation are sensitive to the quality of state information and its uncertainty. The translation from task-graph to execution plan must also trade-off between timely response and near-optimal plans on the one hand, with resilience on the other.

Resilience challenges traditional notions of workflow performance: for example, execution plans must be optimized not for makespan, but for resilient execution. The restatement of the objective is simple, however the objective function of “maximal (optimal) resilience” is a complex, multidimensional function with often ill-defined trade-offs. Providing resilience, say using replication in a way that minimizes correlated failure probabilities, introduces performance costs, such as increased latency or decreased throughput. Workload management must also resolve this tension between resilience and performance.

Workflow Research The workflow middleware will provide a distributed workload management system that can select, acquire and dynamically scale resources across the edge-to-exascale continuum. Specifically, the middleware must enable: (i) resource selection and task placement, and management of these processes so as to allow collective decision making but independent control; (ii) a decision making framework for different modes of data and task placement, while being agnostic to specific resource, services and task properties, and (iii) tunable to diverse performance measures including reliability, response time, power or cost of data movement. These are in addition to associated performance measures such as makespan, latency tolerance, and scale.

Three specific research themes that will be explored: (i) In the presence of heterogeneity and dynamism, fine-grained control and adaptive coordination of the placement of data and tasks is needed. We will explore workload management and associated algorithms to provide the basis for an initial execution plan; (ii) We will investigate robust decision framework (based on models of system behavior) to respond to dynamism in the context of data-task execution and resource topology; (iii) We will perform the system software research to provide the collective and end-to-end integration of capabilities while preserving performance and resilience requirements.

An affinity-based workload management model predicated on the trifecta of data-resource-task, and the coupling of individual data units, resources and tasks (as measured by their relative “affinity” to other units) allows planning task placement, and the flow and processing of data on the edge. The data-resource-task affinity model provides the ability to influence the placement of data-resources-task without exposing either implementation details or “how”. It will support diverse scenarios (e.g., data flows out to the edge, data-

joins from distributed edge locations etc.). A unique feature of Laminar will be the ability to decompose workloads into functions for fine-grained placement (e.g., device level) that can be aggregated to support more coarse-grained placement (e.g., on the edge, private cloud or exascale system). The granularity versus placement trade-off also present powerful adaptive response opportunities.

Because execution plans and predictions do not always agree with reality, we formulate a data-model “discrepancy function” about which we will reason about adaptive response. The discrepancy is an unknown function over which we can express a Bayesian uncertainty to be minimized as a middleware operational objective. However, many open research and implementations questions related to a resilient middleware that supports robust (decision making) adaptations must be addressed: How to formulate an adaptable execution plan and quantify uncertainty in adaptation? How to accurately assess the cost of a change in execution plan and implication for resilience? How to assess when to adapt and how to configure resources while preserving resilience and robustness?

The XGFABRIC system component of distributed resource management system will be the pilot abstraction. It will acquire resources and be responsible for the fine-grained resource partitioning and assignment on the diverse range of resource end-points. The immediate challenge here will be the design of a pilot system that is functionally capable and consistent with the established pilot paradigm but realized on multi-scale resources at very different points in the performance-persistence property space.

We will develop the interfaces to task-level services provided by Laminar and CSPOT needed by the workload management system, e.g., monitoring data, tools for getting data for analysis. To a CSPOT task the pilot system will be accessible as an external service, agnostic of the task semantics and internal CSPOT optimizations. The pilot system will expose well-defined interface that allow CSPOT programming system to probe at runtime to determine state, and declaration of where/when workloads are destined.

Information flow and interface considerations aside, fundamental challenges arising from integrating different system abstractions for distributed resource management (pilot) and runtime systems (CSPOT) need to be addressed. For example, how the pilot bootstraps CSPOT will likely differ across resource types. Two obvious modes warrant investigation: (M1) using the pilot to launch CSPOT tasks, which in turn manages tasks, and/or (M2) Using the pilot as an execution environment running within CSPOT.

In the case of (M1), the hierarchy of XGFABRIC abstraction will be that the glide-in acquires and configures resources, and delivers it to the pilot; the pilot starts its management components, which controls the execution of computing tasks (self-contained processes) on acquired resources; the CSPOT execution environment will in turn be described as a computing task, and will be processed by the pilot as a special task operating within the confines of the pilot.

In the case of M2, the execution chain that the glide-in acquires and configures resources, and launches CSPOT; then an application using CSPOT starts and interfaces with pilot, which controls the execution of tasks. Modes M1 and M2 differ in their resilience vs performance profiles, as well as how they qualitatively and quantitatively respond to dynamism. It is likely that M2 will be initial mode on edge and IoT devices. The challenge will be managing the cross-over behavior to M1 – which is the currently envisioned mode on HPC/data-center/cloud environments.

Evidence that this approach will succeed: Our workflow approach employs the “classic” three-level architecture. The three levels provide a trade-off between complexity and capability, extensibility and specificity, as well as generality and performance.

The pilot abstraction [6, 19, 21] is both powerful and general to employ across the multi-scale resources. To leverage this utility, XGFABRIC research will focus on “right-sizing” the selection of resources to be allocated to each pilot in a multi-scale deployment (i.e. from edge to HPC platform). To guide both qualitative and quantitative insights, we will also develop an emulator to investigate new algorithms and adaptive execution plans. The emulator will hide the complexity of the distributed environment, allow the project to validate new operational models, and derive insight without having to replicate the complex distributed and

dynamic environment. Finding optimal adaptive execution plans – already an NP-hard problem – in the presence of application performance constraints will require heuristics, and solutions with bounded deviations from optimality. As scale and dynamism increases, or resource (task) heterogeneity increase, improvements due to adaptivity will continue to increase. We will use the emulator for instant insight and judicious choice of algorithms, but not be constrained by it.

5 Work Track 3: Deployment, Testing and Validation

5.1 Deployment via Glide-In: The FabriStack as Middleware (UCSB, BNL, Notre Dame)

To demonstrate a usable end-to-end system spanning the entire range from IoT sensors to HPC facility, connected by 5G/6G networks, a global orchestration component must be created. The XGFABRIC orchestrator will accept an application definition, assess the availability of resources at the requested sites, then allocate appropriate resources from the edge, network and HPC facilities, and then deploy the XGFABRIC software. Once the workflow has run to completion, the outputs will be extracted, and the resources cleaned up.

The Research Challenge: The fundamental research challenge is to provide the system software facilities to “speed match” the parts of a streaming application that are immediately available when ever they are presented with data (and probably on some regular duty cycle) with the batch-controlled resources that must be provisioned in priority order so as to maintain good system resource utilization. Further, this alignment cannot mandate that the HPC resources install new software to meet this challenge. To leverage large-scale high-performance computing (HPC) resources as essential computing elements requires that the FabriStack be able to act as a “guest” when using these resources. Put another way, we do not expect that CSPOT, Linux containers, or “root” Linux access will be available to the FabriStack as native technologies from batch-controlled HPC resources. Thus, our approach is to “glide in” FabriStack functionality, dynamically, as middleware that operates as an unprivileged guest on these systems. Addressing this challenge is difficult because HPC resources appear, to applications using the FabriStack, as intermittently available resources. They are reliable (relative to small devices) in that they do not fail often once they become available, but batch queuing delays interpose availability.

Glide In for HPC resources Integrated by XGFABRIC: Our research approach for implementing FabriStack glide in will consist of three ideas for matching the cadence of the sensor or streaming applications outside of the an HPC facility with the “bursty” availability caused by batch queuing. These advances are predicated on research in adaptive resource management and response. The first is to provision staging areas, using the FabriStack, in private or public clouds where resources can be dynamically acquired and released “on demand.” The purpose of these staging areas is to hold the inputs needed by an HPC computation once it has emerged from a batch queue and begins executing.

Note that Laminar’s strict dataflow semantics will not trigger a macro node before all of its inputs are available. Our idea is to introduce a “pass through” node type into Laminar that will be assigned to a cloud (or other resource with sufficient capacity to stage the inputs needed by a large HPC computation). This pass-through node will be inserted into a Laminar program immediately ahead of any macro node assigned to an HPC resource. The pass through node will not “fire” (send its outputs to the HPC macro node) until all of the inputs are present. The Laminar and CSPOT retry mechanisms will then treat the HPC macro node that takes its inputs from the pass through node as if it is available across a potentially faulty network link. The computation corresponding to the macro node will be submitted to the batch queue and we will configure the pass through node to continue to retry the initiation of the macro node computation while the computation waits for execution in the queue.

We will need to implement the ability to instantiate a CSPOT deployment in a cloud dynamically to implement the pass-through node. If private clouds are available, we will dedicate some amount of private cloud resource to HPC staging of Laminar pass-through nodes. However, for public clouds (where idle resources incur unnecessary charges) we will dynamically instantiate CSPOT. We will explore using a pi-

lot [21] to effect the actual glide in. Once the computational elements of the HPC resource are provisioned to the macro node, the pilot will execute the tasks that are necessary to contact the pass-through node, acquire the staged inputs, start the computational tasks, and gather the results.

Evidence that this approach will succeed: Of the three components that comprise the FabriStack, we have the least preliminary experience with Glide In. CSPOT does have a replay facility [15] that is primarily for data repair when sensor data is deemed to be faulty, or when bugs are discovered and repaired in inline computations within a streaming application. We believe we will be able to leverage this capability to dynamically provision CSPOT WooFs within a running Laminar application by creating an empty WooF and then “repairing” it as if it simply contained faulty data. If this approach fails, we have also developed a highly resilient implementation of a replicated “pub-sub” service based on Chord [31] and RAFT [23] (using CSPOT as the runtime) called Canal [16]. Canal is prepared to use RAFT’s strongly consistent replication semantics and Chord’s churn mechanisms to dynamically acquire and release CSPOT hosts. It is not yet clear whether log repair will be sufficient to implement Laminar pass-through nodes or whether a more resilient system such as Canal will be necessary.

5.2 Testbed and Deployment Plan

The XGFABRIC testbed will consist of a **software repository** containing the latest releases of each of system components, a **application repository** containing exemplar software applications that can be deployed, and a set of **experimental facilities** available to the team, noted in the Facilities appendix below.

We expect at the outset of the project that each software effort will quickly produce a “minimally compliant” artifact that can be used as the baseline to allow for integration testing and evaluation of facility and application constraints. As the various software components (SensorSlicer, Laminar, CSPOT, Laminar, Glide-In) develop along their own tracks, the XGFABRIC software repository will identify milestone versions of these components and perform automated tests to evaluate basic internal compatibility. This will result in “known good” configurations that all project participants can rely upon for successful execution. Known good configurations will then be tested against benchmark applications on the available testbed facilities, which consist of the AgIoT facilities (e.g. CUPS at UCSB, Spidercam at UNL) combined with HPC facilities (e.g. HCC at UNL, CRC and ND). As integrating testing reveals inter-component problems, or component-facility problems, these will be fed back to the individual tracks for debugging and refinement. Once the integration produces working combinations, we will proceed to validation and testing.

5.3 Validation Objectives

We will consider the end-to-end execution of XGFABRIC to be **correct and successful** if a user at a “neutral” site is able to orchestrate the deployment, operation, and teardown of an analysis workflow that spans an agricultural facility and an HPC facility, connected by a 5G/6G network, without fixed allocations at either of those sites. For example, an operator physically located at Notre Dame should be able to connect the CUPS facility at UCSB with an HPC facility at BNL, execute the desired workflow, and then fully cleanup.

While the successful coordinated execution of such a complex system is the primary objective, it is also necessary to have sufficient performance to meet application needs. We will observe and address performance in the following dimensions: **Whole-system bandwidth and latency.** The system performance is ultimately constrained by the ability to move data all the way from the sensor, over the wireless network, to the relevant node on the HPC facility. We will measure to what extent the maximum wireless network performance is reduced by the end-to-end system components. **Resource allocation efficiency.** When allocating coordinated resources from multiple facilities (edge, network, HPC), it is unavoidable that some mismatch will result in under-use of one resource or another. (For example, the available network bandwidth may exceed what the edge facility is capable of producing.) We will measure the allocated and used capacity of each resource with an eye towards maximizing utilization withing performance constraints. **End-to-end iteration time.** The ability to take meaningful action on agricultural facilities will be limited by the end-to-

end iteration time needed to deploy, measure, simulate, and perform an actuation on the facility.

6 Timetable of Activities and Milestones

Activities are organized to ensure the rapid integration between research insight, middleware and system development, and science drivers, This is reflected in yearly integration and usage for science milestones (Track 3) based upon XGFABRIC releases 1.0 and 2.0 (Table 1).

The evaluation and impact of individual work packages are discussed in individual work package sections. Project-level evaluation metrics will include: (i) Uptake of prototypes beyond this project will be a consideration; (ii) ease of deployment and portability of applications across diverse test-beds and exascale systems, (iii) resiliency to faults and intermittent connectivity across the wide area, the efficiency of placement/migration decisions and the speed with which decisions can change in response to the system behavior and resource availability. Finally, we will also measure the speed and accuracy with which our prototypes can be controlled and adapted in response real-time data (fused from multiple sources).

To advance scientific goals, XGFABRIC will be deployed and integrated on platforms of practice, DOE leadership-class computing facilities, and other institutional testbeds and resources.

	Month 6	Month 12	Month 18	Month 24
Track 1	<ul style="list-style-type: none"> • Deploy srsRAN and OAI private cells at ENREEC • Extend CSPOT support to exascale systems • Define abstractions and services for geo-distributed monitoring 	<ul style="list-style-type: none"> • Deploy 5G network slicing primitives • Define CSPOT APIs/services for WMS runtime capabilities • Specialize services to digital agriculture workloads and deployments 	<ul style="list-style-type: none"> • Implement and deploy SensorSlicer • APIs/Services for data fusion from disparate sources • Tooling for validation and stress/fault testing 	<ul style="list-style-type: none"> • APIs for adaptive sensing • CSPOT task decomposition support for WMS • End-to-end validation and stress/fault testing
Track 2	<ul style="list-style-type: none"> • Architecture, component and interface design for multiscale workload management system • Affinity model for data-resource-task 	<ul style="list-style-type: none"> • Resource selection and task placement algorithms using affinity model • Implement pilot system for XGFABRIC 	<ul style="list-style-type: none"> • Adaptation using dynamic information and optimization • Bidirectional system integration with CSPOT 	<ul style="list-style-type: none"> • Evaluate and improve heuristics for optimal execution plans • Performance, Scale and Benchmarks
Track 3	<ul style="list-style-type: none"> • Collect minimally complete software components. • Collect science drivers. 	<ul style="list-style-type: none"> • Develop orchestration and test on minimal components and science drivers. • Test integration of 1.0 components on testbeds. 	<ul style="list-style-type: none"> • Develop resilience techniques for variable resource availability. • Test integration of 2.0 components on testbeds. 	<ul style="list-style-type: none"> • Develop techniques for capacity matching between facilities. • Evaluate performance of 3.0 components on testbed.
Milestones	<ul style="list-style-type: none"> • Develop first minimal complete components • Deploy 5G private network 	<ul style="list-style-type: none"> • XGFABRIC 1.0 and initial integration with science driver • Deployed on Exeter test-bed 	<ul style="list-style-type: none"> • XGFABRIC full integration with additional science drivers • XGFABRIC 2.0 	<ul style="list-style-type: none"> • XGFABRIC deployed on Sedgewick, Spidercam test-beds & DOE platforms • XGFABRIC 3.0

Table 1: Timetable of Project Activities and Milestones.

Appendix 1 – References

- [1] OpenAirInterface. <http://www.openairinterface.org/>.
- [2] srsRAN Project. <https://www.srsran.com>.
- [3] Notice of proposed rulemaking. Technical report, FCC, Aug. 2019. <https://docs.fcc.gov/public/attachments/FCC-19-77A1.pdf>.
- [4] Radhakisan Baheti and Helen Gill. Cyber-physical systems. *The impact of control technology*, 12(1):161–166, 2011.
- [5] Pete Beckman, Charlie Catlett, Moinuddin Ahmed, Mohammed Alawad, Linqun Bai, Prasanna Balaprakash, Kevin Barker, Pete Beckman, Randall Berry, Arup Bhuyan, Gordon Brebner, Klaehn Burkes, Anastasiia Butko, Franck Cappello, Ryan Chard, Scott Collis, Johnathan Cree, Dipankar Dasgupta, Anatoly Evdokimov, Jason M Fields, Peter Fuhr, Colby Harper, Yier Jin, Rajkumar Ket-timuthu, Mariam Kiran, Robert Kozma, Praveen Ashok Kumar, Yatish Kumar, Linqing Luo, Lena Mashayekhy, Inder Monga, Bill Nickless, Thrasyvoulos Pappas, Elena Peterson, Trever Pfeffer, Shaloo Rakheja, Veroica Rodriguez Tribaldos, Sterling Rooke, Sumit Roy, Tarek Saadawi, Alec Sandy, Rajesh Sankaran, Nicholas Schwarz, Suhas Somnath, Marius Stan, Cory Stuart, Ryan Sullivan, Anirudha Sumant, Greg Tchilinguirian, Nhan Tran, Arun Veeramany, Angela Wang, Bin Wang, Andrew Wiedlea, Stijn Wielandt, Theresa Windus, Yuxin Wu, Xi Yang, Zhi Yao, Rose Yu, Yuping Zeng, and Yuepeng Zhang. 5g enabled energy innovation: Advanced wireless networks for science. *PNNL Workshop Report*, 3 2020.
- [6] Georgios Chantzalexou, Andre Luckow, and Shantenu Jha. Pilot-streaming: A stream processing framework for high-performance computing. In *2018 IEEE 14th International Conference on e-Science (e-Science)*, pages 177–188. IEEE, 2018.
- [7] X. Dong, M.C. Vuran, and S. Irmak. Autonomous precision agriculture through integration of wireless underground sensor networks with center pivot irrigation systems. *Ad Hoc Networks*, 11(7):1975–1987, 2013.
- [8] Andy Rosales Elias, Nevena Golubovic, Chandra Krintz, and Rich Wolski. Where’s the bear?-automating wildlife image processing using iot and edge cloud systems. In *2017 IEEE/ACM Second International Conference on Internet-of-Things Design and Implementation (IoTDI)*, pages 247–258. IEEE, 2017.
- [9] John Ellson, Emden R Gansner, Eleftherios Koutsofios, Stephen C North, and Gordon Woodhull. Graphviz and dynagraph—static and dynamic graph drawing tools. *Graph drawing software*, pages 127–148, 2004.
- [10] UCSB Early Research Scholars Program, 2023. <https://ersp.cs.ucsb.edu> [Online; accessed 3-Feb-2023].
- [11] Espressif 8266 microcontroller. <https://en.wikipedia.org/wiki/ESP8266>, 2018. [Accessed electronically, September 2018].
- [12] EM Fajardo, JM Dost, B Holzman, T Tannenbaum, J Letts, A Tiradani, B Bockelman, J Frey, and D Mason. How much higher can htcondor fly? *Journal of Physics: Conference Series*, 664(6), 2015.
- [13] 2024. <https://farm-ng.com> [Online; accessed 5-Apr-2024].

- [14] Nevena Golubovic, Rich Wolski, Chandra Krintz, and Markus Mock. Improving the accuracy of outdoor temperature prediction by iot devices. In *2019 IEEE International Congress on Internet of Things (ICIOT)*, pages 117–124. IEEE, 2019.
- [15] W-T. Lin, F. Bakir, C. Krintz, R. Wolski, and M. Mock. Data repair for Distributed, Event-based IoT Applications. In *ACM International Conference On Distributed and Event-Based Systems*, 2019.
- [16] Wei-Tsung Lin, Rich Wolski, and Chandra Krintz. A programmable and reliable publish/subscribe system for multi-tier iot. In *2021 8th International Conference on Internet of Things: Systems, Management and Security (IOTSMS)*, pages 1–8. IEEE, 2021.
- [17] K. LoPiccalo. Impact of broadband penetration on U.S. farm productivity. Technical report, Federal Communications Commission, Office of Economics and Analytics, Feb. 2021. <https://docs.fcc.gov/public/attachments/DOC-368773A1.pdf>.
- [18] University of California Division of Agricultural and Natural Resources, Lindcove Research and Extension Center. <http://lrec.ucanr.edu>. [Online; accessed 13-mar-2023].
- [19] Andre Luckow, Kartik Rattan, and Shantenu Jha. Pilot-edge: Distributed resource management along the edge-to-cloud continuum. *arXiv preprint arXiv:2104.03374*, 2021. Accepted for PAISE’21 (IPDPS 21).
- [20] Mohammad M. R. Lunar, Jianxin Sun, John Wensowitch, Michael Fay, Halit Bugra Tulay, Sai Suman, Brian Qiu, Deepak Nadig, Garhan Attebury, Hongfeng Yu, Joseph Camp, Can Emre Koksall, Dario Pompili, Byrav Ramamurthy, Morteza Hashemi, Eylem Ekici, and Mehmet C. Vuran. OneLNK: one link to rule them all: Web-based wireless experimentation for multi-vendor remotely accessible indoor/outdoor testbeds. In *ACM Workshop on Wireless Network Testbeds, Experimental evaluation and Characterization (WiNTECH’21)*, Feb. 2021.
- [21] Andre Merzky, Matteo Turilli, Mikhail Titov, Aymen Al-Saadi, and Shantenu Jha. Design and performance characterization of radical-pilot on leadership-class platforms. *IEEE Transactions on Parallel and Distributed Systems*, 33(4):818–829, 2021.
- [22] James Murty. *Programming amazon web services: S3, EC2, SQS, FPS, and SimpleDB*. ” O’Reilly Media, Inc.”, 2008.
- [23] Diego Ongaro and John Ousterhout. In search of an understandable consensus algorithm. In *2014 {USENIX} Annual Technical Conference ({USENIX}{ATC} 14)*, pages 305–319, 2014.
- [24] Philippe Rolshausen. Prospects for farming citrus under protective screening. *Citrograph*, 14(4), 2023.
- [25] Mendel Rosenblum and John K Ousterhout. The design and implementation of a log-structured file system. *ACM Transactions on Computer Systems (TOCS)*, 10(1):26–52, 1992.
- [26] Abdul Salam, Mehmet C. Vuran, Xin Dong, Christos Argyropoulos, and Suat Irmak. A theoretical model of underground dipole antennas for communications in internet of underground things. *IEEE Transactions on Antennas and Propagation*, 67(6):3996–4009, jun 2019.
- [27] A.R. Silva and M.C. Vuran. (CPS)²: Integration of center pivot systems with wireless underground sensor networks for autonomous precision agriculture. In *Proceedings of the 1st ACM/IEEE International Conference on Cyber-Physical Systems, ICCPS ’10*, pages 79–88, 2010.

- [28] A.R. Silva and M.C. Vuran. Development of a testbed for wireless underground sensor networks. *Eurasip Journal on Wireless Communications and Networking*, 2010, 2010.
- [29] Stephen K Skedzielewski and ML Welcome. Data flow graph optimization in if1. In *Functional Programming Languages and Computer Architecture: Nancy, France, September 16–19, 1985*, pages 17–34. Springer, 1985.
- [30] Stephen K Skedzielewski and ML Welcome. Data flow graph optimization in if1. In *Functional Programming Languages and Computer Architecture: Nancy, France, September 16–19, 1985*, pages 17–34. Springer, 1985.
- [31] Ion Stoica, Robert Morris, David Liben-Nowell, David R. Karger, M. Frans Kaashoek, Frank Dabek, and Hari Balakrishnan. Chord: a scalable peer-to-peer lookup protocol for internet applications. *IEEE/ACM Trans. Netw.*, 11(1):17–32, 2003.
- [32] Fei Tao, He Zhang, Ang Liu, and Andrew YC Nee. Digital twin in industry: State-of-the-art. *IEEE Transactions on industrial informatics*, 15(4):2405–2415, 2018.
- [33] J. Tooker, X. Dong, M.C. Vuran, and S. Irmak. Connecting soil to the cloud: A wireless underground sensor network testbed. In *Annual IEEE Communications Society Conference on Sensor, Mesh and Ad Hoc Communications and Networks workshops*, volume 1, pages 79–81, 2012.
- [34] N. M. Trendov, S. Varas, and M. Zeng. Digital technologies in agriculture and rural areas – status report. Technical report, Food and Agriculture Organization of the United Nations, Rome, Italy, 2009. Licence: cc by-nc-sa 3.0 igo.
- [35] UC Agriculture and Natural Resources, 2023. <https://ucanr.edu>[Online; accessed 3-Feb-2023].
- [36] UC Natural Reserve System, 2023. <https://ucnrs.org>[Online; accessed 3-Feb-2023].
- [37] UCSB RACELab, 2023. <https://sites.cs.ucsb.edu/~ckrintz/racelab.html>[Online; accessed 3-Feb-2023].
- [38] 2019. <http://sedgwick.nrs.ucsb.edu> [Online; accessed 14-Feb-2019].
- [39] Mehmet C. Vuran, Abdul Salam, Rigoberto Wong, and Suat Irmak. Internet of underground things in precision agriculture: Architecture and technology aspects. *Ad Hoc Networks*, 81:160–173, dec 2018.
- [40] Rich Wolski, Chandra Krintz, Fatih Bakir, Gareth George, and Wei-Tsung Lin. Cspot: portable, multi-scale functions-as-a-service for iot. In *Proceedings of the 4th ACM/IEEE Symposium on Edge Computing*, page 236–249. ACM, Nov 2019.
- [41] Zhongyuan Zhao, Mehmet C. Vuran, Baofeng Zhou, Mohammad M.R. Lunar, Zahra Aref, David P. Young, Warren Humphrey, Steve Goddard, Garhan Attebury, and Blake France. A city-wide experimental testbed for the next generation wireless networks. *Ad Hoc Networks*, 111:102305, Feb. 2021.

Appendix 2 – Facilities

Brookhaven National Lab

Computational Science Initiative

While long focused on timely analysis, interpretation, and overall management of high-volume, high-velocity heterogeneous data to pursue solutions for the national and international scientific community, the Computational Science Initiative (CSI) at the U.S. Department of Energy's (DOE) Brookhaven National Laboratory additionally excels at integrating computer science, applied mathematics, and computational science with broad domain science expertise to tackle problems and advance knowledge impacting scientific discovery. CSI's expertise and investments across the Brookhaven Lab, including its connectivity to flagship physics and materials science facilities that attract thousands of scientific users each year, are tackling the most pressing big data and science challenges. These efforts now are being augmented by CSI's growing high-performance computing (HPC), applied mathematics, artificial intelligence/machine learning, and quantum science capabilities.

Scientific Data and Computing Center

Brookhaven's Scientific Data and Computing Center (SDCC) combines advanced expertise in high-throughput, high-performance, and data-intensive computing with data management and preservation in a centralized computing facility. It provides varied services to local and international clients that require stable, reliable computing capabilities for processing, storing, and analyzing large-scale data sets, along with HPC resources for increased computing power.

The SDCC houses systems for high-performance and data-intensive computing, data storage, and networking, offering everything from novel research platforms to highly reliable production services. The facility stores 230 PB of data. In 2022, the SDCC processed 1.4 exabytes of data and transferred in and out more than 200 PB of data, ranking it among the top 10 data archives in the world.

SDCC High-performance Computing Clusters

Institutional Cluster Gen2 (ICGen2)

The ICGen2 cluster consists of:

- 39 CPU worker nodes
- 12 4xA100-SXM4 GPU nodes
- One 2xA100 80 GB PCIe node
- One 2xA100 40 GB PCIe node
- One submit node
- Two master nodes

CPU nodes:

- Supermicro SYS-610C-TR
- Intel® Xeon® Gold 6336Y CPU @ 2.40 GHz
- NUMA node0 CPU(s): 0-23
- NUMA node1 CPU(s): 24-47
- Thread(s) per core: 1
- Core(s) per socket: 24
- Socket(s): 2
- NUMA node(s): 2
- 512 GB Memory
- InfiniBand NDR200 connectivity

GPU A100-SXM4 nodes:

- Supermicro SYS-220GQ-TNAR+
- Intel® Xeon® Gold 6336Y CPU @ 2.40 GHz

- NUMA node0 CPU(s): 0-23
- NUMA node1 CPU(s): 24-47
- Thread(s) per core: 1
- Core(s) per socket: 24
- Socket(s): 2
- NUMA node(s): 2
- 1 TB Memory
- 4x A100-SXM4-80 GB
- InfiniBand NDR200 connectivity

GPU 2xA100 (80/40) GB PCIe nodes (debug partition)

- Supermicro SYS-120GQ-TNRT
- Intel® Xeon® Gold 6336Y CPU @ 2.4 0GHz
- NUMA node0 CPU(s): 0-23
- NUMA node1 CPU(s): 24-47
- Thread(s) per core: 1
- Core(s) per socket: 24
- Socket(s): 2
- NUMA node(s): 2
- 512 GB Memory
- 2xA100 (80/40) GB (amperehost01/amperehost02)
- InfiniBand NDR200 connectivity

Storage:

- 1.9 TB of local disk storage per node
- 1 PB of GPFS distributed storage
- Cluster Storage

NSLS2 Cluster

- 32 Supermicro nodes with EDR IB
- 13 nodes with 2x NVIDIA V100

SDCC High-throughput Computing (HTC)

- ~ 1,900 HTC nodes available to SDCC users:
 - ~90,000 logical cores
 - ~1050 kHS06
 - Managed by HTCondor 9.0
 - All nodes running Scientific Linux (SL) 7

IC Central Storage

BNLBox is SDCC's cloud storage service that enables users to store their scientific data and documents locally, providing accessibility that is available from anywhere in the world via Internet access. The files stored in BNLBox can be shared with other BNLBox users, as well as external collaborators.

Computational Sciences

The Computational Sciences Department, also part of CSI, operates a collaborative laboratory for advanced algorithm development and optimization. The focus is to develop tools and techniques to solve problems in computational physics, biology, chemistry, materials science, and energy and environmental sciences by effectively using increasingly faster supercomputers.

Field Programmable Gate Array Development Server

FPGA development server is a Linux server with:

- 1 Intel Xeon CPU E5-2620 with eight cores
- 256 GB Memory
- One Intel Altera 10 GX FPGA card
- One NVIDIA Pascal 100 GPU
- Two NVIDIA Volta 100 GPUs

Brookhaven Lab SciServer Implementation

Johns Hopkins University (JHU)'s SciServer is a fully integrated cyberinfrastructure system with related tools and services (<https://www.idies.jhu.edu/what-we-offer/sciserver/>) to enable researchers to work on tera- or peta-bytes of scientific big data. Hosted and operated by Brookhaven Lab, the Brookhaven SciServer implementation includes a data hosting and analysis infrastructure with an interface for uploading and managing datasets, as well as tools to launch and monitor jobs. BNL's SciServer includes four powerful compute nodes, each with eight NVIDIA H100 GPUs for training artificial intelligence models. These nodes will be connected to multiple petabytes of storage, facilitating collaboration in a federated environment. In addition, BNL and JHU have worked together to integrate key SciServer extensions.

BNL SciServer Implementation

- Four Machine Learning (ML) Compute Nodes. Each includes eight NVIDIA H100 GPUs with SXM form factor, 1.5 TB RAM, and 2x 100 GigE links.
- Five ML Compute Nodes. Each includes eight NVIDIA V100 GPUs and 10 GigE links.
- Thirty-six ML Compute Nodes. Each includes two NVIDIA P100 GPUs and 10 GigE links.
- Five-petabyte Storage.
- Kubernetes cluster and various support servers.

Systems, Architectures, and Emerging Technologies

Part of Brookhaven's CSI, the Systems, Architectures, and Emerging Technologies Department is at the crossroads of research conducted to understand how architectures and systems can tackle the challenges posed by data-intensive computing and its impacts on both scientific and national-security-motivated problems. The department also oversees the Computing for National Security Center.

Advanced Computing Lab (ACL)

The ACL is a focal point and collaborative environment at Brookhaven Lab for research and development (R&D) of architectures and systems, as well as for siting and enabling advanced computing testbeds. The ACL consists of state-of-the-art equipment capable of facilitating R&D on computing technologies at different development and maturity stages, from materials to devices to subsystems to early full-system prototypes. High-resolution, accurate static and dynamic instrumentation for performance (time), power, and reliability can be included on specialized workbenches. The ACL features machine room space properly designed for optimal operation and safe access to experimental testbeds. Network connectivity for external users will be provided as required. The ACL can include office space for technology providers and external collaborators.

NVIDIA DGX-2™ Artificial Intelligence Cluster

A state-of-the-art NVIDIA DGX-2 is fully operational in the Advanced Computing Lab (ACL). The system is powered by two Intel Xeon Platinum CPUs and 16 fully interconnected Tesla® Tensor Core V100 graphics processing units (GPUs), each with 32 GB of HBM2 memory. Total system memory is 1.5 TB of DDR4. Storage is 30 TB of non-volatile storage class memory. The system is connected to Brookhaven's Scientific Data and Computing Center (SDCC) via dual 100 GigE networks. The system software consists of the original operating system, a "singularity" container, and Slurm scheduler. Brookhaven Lab was among the earliest adopters of the system, nicknamed Minerva.

NVIDIA DGX A100

- One 4xH100 GPU server
- Two 8xA6000 GPU servers
- 1 PB Storage

Quantum Network Testbed

With the help of DOE's Energy Sciences Network (ESnet), Brookhaven Lab has built a quantum network testbed on its campus using existing fiber infrastructure. The testbed includes a central hub in the main data center (Scientific Data and Computing Center) connected to the Brookhaven's primary networking hub. From there, entangled photons can be distributed to two on-site quantum technology laboratories for short- and medium-distance measurements

Energy Sciences Network (ESnet)

Brookhaven Lab also is connected to DOE's Energy Sciences Network via a 300 Gb/s link. ESnet provides advanced services to more than 40 DOE research sites, including the entire national laboratory system, its supercomputing facilities, and major scientific instruments. ESnet also connects to 140 research and commercial networks, permitting DOE-funded scientists to collaborate with partners around the world.

UC Santa Barbara

In addition to the UCSB RACELab [37], the project will have access to the Sedgwick Reserve [38] and the Lindcove Research Extension Center [18]. Sedgwick is operated by UCSB as part of the University of California Natural Research System (NRS) [36]. UCSB administers other NRS sites that are potentially available to the project as well, however the existing testbed facilities at Sedgwick are immediately available as a result on an on-going collaboration between the project team members. Lindcove is an agricultural research facility operated by University of California Agriculture and Natural Resources (ANR) [35]. Like the NRS, UCANR operates a number of research stations across California that are potentially available to the project. The project team has an existing testbed deployment at Lindcove (which includes the Flux Tower described in Subsection 5.2 of the Research Narrative) and an on-going collaboration that this project will leverage.

University of Notre Dame

The Notre Dame Center for Research Computing (CRC) operates a state-of-the-art high performance computing infrastructure providing advanced computing support to researchers and teachers at Notre Dame and within the local community. The CRC systems have a wide range of software applications, supporting research across campus. The CRC has extensive experience in software development and performance profiling with numerous engineers and professional programmers on staff. The CRC data center is managed in partnership with Global Access Point's commercial data center in downtown South Bend, IN. This ensures our data center is operated according to industry standards and allows for flexible growth in physical footprint and utility demands. ND CRC system engineers have full time offices in the facility to provide rapid response. Union Station facilities feature three separate power grids with UPS and APS sources, FM-200 fire suppression, advanced security, redundant climate control technology, and advanced sensors monitoring temperature, humidity and smoke. GAP's 24/7/365 Network Operation Center ensures secure customer access and emergency response capabilities.

University of Nebraska-Lincoln

Cyber-Physical Networking (CPN) Laboratory: The 1,300 sq. ft. CPN Lab houses graduate student office space and equipment.

CPN Underground Wireless Sandbox Testbed: An indoor sandbox testbed was developed inside a greenhouse at UNL east campus in the form of a 100"x26"x48" wooden box holding approximately 90ft³ of soil for controlled wireless underground experiments.

Holland Computing Center (HCC): HCC provides access to Nebraska's supercomputing resources at the Peter Kiewit Institute (PKI) in Omaha and Schorr Center in Lincoln. HCC will support by meeting its processing needs, Gbps fiber optic link and have 100 Gbps primary paths with 10 Gbps backup paths to the University of Nebraska, Internet2, and ESnet. A 2,200 sq. ft. machine room at Schorr provides up to 100 ton cooling with up to 400 kVA of power. Dell S4248FB-ON edge switches and Z9264F-ON core switches

provide high WAN bandwidth and Software Defined Networking capability for both locations. Schorr has space, power, and cooling for 10 racks (5-10kW per rack).

HCC is home to the following resources. 1) Crane, which debuted on the Top500 Supercomputer list, contains Intel Xeon 8-core 2.6GHz and 18-core 2.3GHz processors totaling 12,236 cores. 2) Rhino offers 7,040 cores interconnected with Mellanox QDR Infiniband along with 360TB of BeeGFS storage. Each compute node is an R815 server with at least 256 GB RAM and four Opteron 6272 (2.1 GHz) processors. 3) Red has 12,512 job slots in 386 hosts interconnected by a mixture of 1, 10, and 40 Gb ethernet. Red has over 11 PB of storage, is integrated with the Open Science Grid. 4) Anvil, an OpenStack cloud machine, provides an alternative to standard Linux batch systems when needed. Anvil consists of 1,520 cores and 1.5PB of CEPH storage, all connected by 10 Gb networking. 5) Attic and Silo form a near line archive with 1.0 PB of usable storage. Attic is located at PKI, and Silo acts as an online backup in Lincoln. Attic and Silo are connected with 10 Gbps network connections. 6) Common, a high-performance storage array totaling 1.9PB capacity, is available from all HCC resources.

Scott Engineering Center (SEC) Data Center: The 3,110 sq. ft. SEC Data Center at UNL serves as the NEXTT testbed core site. The facility is equipped with 15 core networking cabinets and 60 server cabinets that house central IT and departmental systems. Open space exists for an additional 40 cabinets. Backup power is provided by two independent generator systems (400 kWh and 500 kWh generators), and there are two independent 160 KVA UPS systems. Cabinet power configuration allows for two sources to each machine. The redundant source is public power backed up by UPS and a generator. A raised floor system is used for forced-air cooling with 45 tons of cooling supplied by Liebert CRAC units. There are three different sources of cooling redundancy: campus chilled water, independent glycol loop, and forced air units. Access to Internet2 traffic is available.

NEXTT Indoor Testbed: It is a fully functional indoor testbed that resides inside the CPN lab.

The Eastern Nebraska Research Extension and Education Center (ENREEC): The UNL Eastern Nebraska Research and Extension Center (ENREEC) facility comprises over 9,000 acres of crop, range-land, and animal production, just north of the main campus in Lincoln Nebraska, which utilizes a variety of modern agricultural field equipment. About half of the land base is in row crops. Fifty percent of those row crops are irrigated. The balance of the site is predominately in cool and warm season pastures. The Agronomy and Horticulture Department has 1,000 acres of land at this location that is utilized for studies related to crop production and field crops involving plant breeding, genetics and molecular physiology, plant physiology and production ecology, soil and water science, and weed science. Technologies such as geospatial and precision technologies are being utilized. The ENREEC Farm Operations Group's mission is to help facilitate research and education programs by providing land, equipment, labor, expertise, and services to departments when the resource is not available within the academic department. The farm provides a working laboratory using modern production practices and equipment sized to today's production agriculture. The ENREEC facility owns and operates each piece of agricultural equipment required to perform field operations for row-crop corn and soybean production. Each piece of equipment is connected in real-time via cellular networks to the myjohndeere.com operations center platform. Figure 2 shows the four critical pieces of equipment that will be used in conjunction with our study, a planter, nitrogen applicator, sprayer, and harvester. The ENREEC facility maintains an account to utilize these datasets when necessary.

NU-Spidercam: High Throughput Field Phenotyping Facility: UNL's NU-Spidercam field facility is a cable robot based system built on a 1.0-ac field. The system is designed to collect high resolution and high accuracy field phenotyping data in a semi-controlled field environment. The unique features of the system are as follows.

- A sensor platform is tethered by eight cables via four poles at the corners of the field. The motion and position (XYZ) of the sensor platform are controlled with a repeatability of ± 5 cm.
- Six sensing modules are currently onboard the sensor platform for plant sensing: a high-resolution

RGB camera, a four-band multispectral camera, a thermal infrared camera, a VNIR hyperspectral camera (from 400 to 1000 nm), a portable VNIR spectrometer attached to up-looking and down-looking fiber optical cables, and a 3D LiDAR scanner. They are used to collect phenotyping data from the field plots.

- The field is divided into 128 six-row zones. Each zone is 4.5x6 m in size. A subsurface drip irrigation system is installed in the field, which allows a controlled amount of water to be delivered to individual plots. Controlled amount of nutrient application is also possible through fertigation. This variable rate irrigation system provides flexibility and confidence to design field experiments related to water and nutrient treatments.
- An onsite weather station is established to record micrometeorological variables including total solar and PAR radiation, air temperature, air humidity, wind speed, and precipitation at one-hour interval.

The animal production portion of the Eastern Nebraska Research and Extension Center is comprised of a Specific Pathogen Free (SPF), farrow-to-finish swine operation (3,500 head). Also there is a 2,500 head research feedlot, and a cow-calf unit that manages 3 different cow herds that totals 550 cows. Unique to the center is the ability to conduct integrated cropping and beef production research. During recent years, both the cow-calf unit and the feedlot have begun to incorporate more precision livestock technology. Lack of connectivity in many areas of the center hinders more incorporation of more of this technology. Similarly, precision livestock technology is being applied in the swine unit. Currently, efforts are underway to develop a 2,000 head Feedlot Innovation Center that will contain outside pens as well as pens in a deep-pit covered confinement facility. The goal of the innovation center is to be provide opportunities to test new precision livestock technologies.

Appendix 3 – Equipment

Brookhaven National Lab

The CSI Visualization Laboratory is used for collaborative visualization, as well as visualization of large data sets. CSI's Visualization Laboratory features:

- Total display resolution of nearly 50 megapixels
- Mosaic software and QuadroSync hardware assures seamless interactivity within the multiple display system
- Computer system
 - Three NVIDIA Quadro P5000s GPUs driving the display system and HTC Vive VR headsets
 - 256 GB RAM

In addition to the graphics capability, the system can be used for GPU computing, providing up to 27 teraflops of single-precision peak computing power.

SDCC Institutional Cluster

The SDCC currently operates the institutional cluster (IC) at Brookhaven Lab. The IC clusters include:

Annie Cluster

- 216 compute nodes with:
 - HPE ProLiant XL190r Gen9
 - Two CPUs Intel Xeon® CPU E5-2695 v4 @ 2.10 GHz
 - Non-uniform Memory Access (NUMA) node0 CPU(s): 0-8,18-26
 - NUMA node1 CPU(s): 9-17,27-35
 - Thread(s) per core: One
 - Core(s) per socket: 18
 - Socket(s): Two
 - NUMA node(s): Two
 - 2x NVIDIA K80 (108 compute nodes) or 2x P100 (108 compute nodes) per node (4 K80 or two P100 devices per node)
 - 256 GB Memory
 - InfiniBand EDR connectivity
- Two submit nodes
- Two master nodes

Central storage for the IC is provided by IBM's Spectrum scale file system (GPFS), a high-performance, clustered file system that supports full POSIX semantics.

- 1.9 TB of local disk storage per node
- 1 PB of GPFS distributed storage
- A GPFS-based storage system with a bandwidth of up to 24 GB/s is connected to the IC.

Francis (Knights Landing) Cluster

- 142 compute nodes
 - KOI S7200AP
 - One Intel® Xeon Phi™ CPU 7230 @ 1.30 GHz
 - NUMA node0 CPU(s): 0-255
 - Thread(s) per core: Four
 - Core(s) per socket: 64
 - Socket(s): One
 - NUMA node(s): One
 - 192 GB Memory
 - Dual Rail Omni-Path (Gen1) connectivity

- Two submit nodes

- Two master nodes

Skylake (Knights Landing) Cluster

The cluster consists of 64 compute nodes with:

- 142 Dell PowerEdge R640
- 142 Two CPUs Intel® Xeon® Gold 6150 CPU @ 2.70 GHz
- 142 NUMA node0 CPU(s): 0,2,4,6,8,10,12,14,16,18,20,22,24,26,28,30,32,34
- 142 NUMA node1 CPU(s): 1,3,5,7,9,11,13,15,17,19,21,23,25,27,29,31,33,35
- 142 Thread(s) per core: One
- 142 Core(s) per socket: 18
- 142 Socket(s): Two
- 142 NUMA node(s): Two
- 142 192 GB Memory
- 142 InfiniBand EDR connectivity

HPC1 Cluster

The HPC1 cluster consists of 21 nodes with 336 cores, each with 128 GB dynamic random-access memory (DRAM), connected with FDR InfiniBand. HPC1 Includes 12 NVIDIA GPUs and 8 Intel Phi coprocessors with 0.5 PB Lustre storage. Field Programmable Gate Array Development Server FPGA development server is a Linux server with:

- 1 Intel Xeon CPU E5-2620 with eight cores
- 256 GB Memory
- One Intel Altera 10 GX FPGA card
- One NVIDIA Pascal 100 GPU
- Two NVIDIA Volta 100 GPUs

NVIDIA DGX-2™ Artificial Intelligence Cluster

A state-of-the-art NVIDIA DGX-2 is fully operational in the Advanced Computing Lab (ACL). The system is powered by two Intel Xeon Platinum CPUs and 16 fully interconnected Tesla® Tensor Core V100 graphics processing units (GPUs), each with 32 GB of HBM2 memory. Total system memory is 1.5 TB of DDR4. Storage is 30 TB of non-volatile storage class memory. The system is connected to Brookhaven’s Scientific Data and Computing Center (SDCC) via dual 100 GigE networks. The system software consists of the original operating system, a “singularity” container, and Slurm scheduler. Brookhaven Lab was among the earliest adopters of the system, nicknamed Minerva.

UC Santa Barbara

The RACELab, at UCSB, operates a number of edge clouds consisting of between 6 and 16 Intel NUC devices (3.5 GHz), with 32 GB of memory, and 1TB of disk each. The clouds run the Eucalyptus open source private cloud software which makes each appear as if it is a stand-alone AWS cloud. Thus, each edge cloud is a miniturized version of AWS. The cloud used a converged storage approach (based on Ceph) to create 3 replicas of all object store data (in the Eucalyptus implementation of S3) and volume storage.

These research labs as part of past work on the NSF UCSB SmartFarm project (NSF CCF15-39586), share an sensor network of microcontroller devices, minicomputers (raspberry Pi and Intel NUC systems), and various sensor devices (for soil moisture, weather, energy use, mobile equipment GPS, and motion trigger image capture and measurement). These devices number in the range of 50 sensors, 10 minicomputers, and 50 sensing devices, all of which will be available for use by the project described herein.

The research team will also have access to the UCSB campus cloud (which also runs Eucalyptus) that is part of the NSF Aristotle project (NSF ACI-1541215 now maintained as campus infrastructure at UCSB). Aristotle is a multi-institution collaboration between UCSB, Cornell, and The University of Buffalo. Together they have developed (and continue to develop) a multi-campus cloud federation, which currently supports seven science teams across sites in astronomy, climatology, economics, biology, sustainable food production and water use, and geospatial data analytics. The UCSB Aristotle cloud consists of approximately 30 nodes (288 cores) and 50TBs of storage. The cloud supports “bursting” (the temporary use of spare, unused resources) to other clouds in the system.

Additionally, the project team will be able to make use of the Knot and Pod clusters in the California NanoSystems Institute (CNSI) at UCSB. The Knot cluster is managed as a batch system with supporting approximately 1500 cores in 120 nodes and 12 GPUs. The Pod cluster supports approximately 2600 cores in 70 nodes interconnected by OmniPath.

We will also work closely with Extension Specialists for the Univ. of California Agriculture and Natural Resources (UC ANR) at the Lindcove Research Extension Center (LREC: <http://lrec.ucanr.edu>). LREC will provide our team with access to sensor-instrumented citrus field blocks, a “pack line,” fruit chemical and physical analysis machine located at LREC. The pack line measures fruit size, number, color, weight, shape, internal and external defects, brix (sugar content), dry matter, and other parameters. The laboratory measures the sugar content, titrated acids, rind thickness, color, and external and internal texture of fruit. LREC and will house multiple edge computing systems. The team also operates a Flux Tower, located at LREC, in Subsection 5.2 of the research narrative.

Thus LREC will provide the team with a production agricultural environment the deployment, validation, testing, and demonstration of the team’s carbon footprint advances and applications.

University of Notre Dame

The Notre Dame CRC provides the following services and resources in support of research. Typically HPC systems oriented toward batch job submission. The CRC manages 20,000+ CPU cores in systems of various architectures and interconnects. Computational resources typically have associated disk systems for short term storage. Storage resources approximately 3PB of data storage including disk based systems for high performance (Panasas, Ceph) and user space storage (NFS, AFS) and tape based systems for backup storage.

The CRC also maintains visualization systems, systems for virtual hosting, prototype architectures, and high throughput computing infrastructure.

In addition to the University-level resources of the CRC, PI Thain operates two shared facilities:

- The Data Intensive Science Cluster (DISC-2), a storage accelerator cluster for caching and managing hot data, consisting of 4TB RAM, 4.8TB NVMe, and 168TB Disk. The storage layer of this cluster is organized as a Ceph parallel filesystem, which is then made accessible to outside users through a range of user-level I/O interfaces, including Chirp, MinIO (S3), XRootD, and HTTP. the DISC cluster provides a location for prototyping and evaluating custom I/O services alongside production computing systems. Internal funding for this cluster was awarded in Fall 2022 and will be procured in Spring 2023.
- A 10K-core HTCondor pool, drawn from computing resources around the University, including the Computer Science and Engineering department, and the Center for Research Computing. This facility allows for opportunistic scaling and testing of applications at the 1000-node scale prior to running reserved submissions on the CRC.

University of Nebraska-Lincoln

Cyber-Physical Networking (CPN) Laboratory: The 1,300 sq. ft. CPN Lab, directed by PI Vuran, includes two Telegraph TGSounder mmWave systems, Severs mmWave front-end phased-array antenna systems, seven USRP-N310, four USRP-N210, seven USRP-B200, two USRP-B210, two USRP-E312, and two BladeRF software-defined radios as well as over 20 low-cost software-defined radio receivers, which can be used for spectrum sensing, cognitive-radio, and vehicle-to-barrier communication experiments. The CPN testbed supports remote programming, out-of-band monitoring, power management, and virtual sensing. The lab includes equipment for assembling and testing hardware platforms and prototypes and equipment for soldering and assembling SMT and through-hole devices, such as a soldering station and hot plate. Assembled devices can be tested and debugged with a variety of test equipment, including a multimeter, oscilloscope, logic analyzer, and spectrum analyzer. Through their use, the lab produces high-quality devices and ensures their correct operation prior to deployment.

CPN Underground Wireless Sandbox Testbed: The sandbox testbed contains 433MHz dipole, 915MHz dipole, GSM quadband, and CPN wideband planar antennas at depths of 10, 20, 30,40cm. Each set of antennas are separated by 50cm from the adjacent sets. To provide cloud-native remote wireless experimentation capabilities to the sandbox testbed, each SDR and IoT user-end device is equipped with a RaspberryPi 4 compute node connected to the sandbox framework via WiFi. Each of these compute nodes is also equipped with LTE to enable out-of-band management. To enable remote user access and experimentation on the UG sandbox, remote desktop to each of the UG nodes, both IoT and SDR, has been integrated into the web framework. Google chrome remote desktop API is used to provide a registered NEXTT sandbox user access to each of the nodes. The users have the capability to create the experiments on the UG SDR sandbox devices using GNURadio. Once the experiment begins, the user may log off until the experiment is finished. UG experiments are conducted over time, usually a few days or weeks, to capture the moisture changes in the sandbox accurately. While controlled experiments are performed at the indoor sandbox testbed, an urban outdoor testbed located conveniently just outside the greenhouse, helps with experiments in real soil conditions.

NEXTT Indoor Testbed: This testbed consists of radio devices from three different vendors. It has two USRP N310 from NI and two mmWave channel sounders from Telefar. A pair of each of these devices are placed in the position and height-adjustable pegboards. These pegboards are installed on two sidewalls of the CPN lab having 49 feet of separation. As a result, both transmitter and receiver radios can be simultaneously configured and controlled by the testbed users. In addition, 4 x 1 linear antenna arrays are installed with each USRP N310 for providing the MIMO experimentation facility. All these radio devices are connected with

a network switch which directly linked with the HCC network through a dedicated fiber connection. This switch provides full remote operability for all radio devices in the testbed through HCC VPN. Necessary compute and storage nodes can also be accessed from this switch. Pegboards are installed on the east side and west side walls of the CPN lab. Antennas of NEXTT radios can be positioned up to 8.5 feet high in this pegboard.

NEXTT Site at ENREEC: ENREEC agronomy building houses a USRP N310 with compute cluster. A 2 x 2 MIMO DVB-T tv antenna is installed on both Tx and Rx frontend of USRP N310. Moreover, a PCTEL GPS antenna is also with it. All these antennas are installed with lightning arrestors so that USRP N310 is protected from lightning. An advanced HPC Mercury 408 4U Tower ServerAH-GPU408-SB01 with Tesla V100-PCI-E-16GB GPU is also on this site for providing real-time high-performance computing. Both compute cluster and USRP are connected with a network switch and fed by a remote-controlled power distribution unit (PDU) unit. Using this PDU these devices can be power cycled remotely. The overall network is connected with HCC VPN. As a result, experimenters can completely configure, execute their experiments remotely. In addition, experiment data can be collected to a preferred storage or cloud server using ENREC's fiber-optic connection. For remote management and health monitoring of all these devices, a cellular gateway from Opendgear is also available on this site. This cellular gateway ensures out-of-band (OOB) connectivity for the system.

Appendix 4 – Data Management Plan

Coordination of Data Management. As the lead institution, BNL will have the primary responsibility for the management of data. Work undertaken at USCB, UNL and Notre Dame resulting in data and software products will generally be incorporated into the primary repositories established at BNL as it is completed. At BNL, all data—software, publications, scientific results and curricula material—will be managed within latest Laboratory and Federal Government guidelines. BNL will help support and enforce these guidelines as per Lab policy. Likewise at Universities (Notre Dame, USB and UNL), all software and data will be managed according to local institutional principles while active, and then will be transferred to BNL facilities for archival and sharing according to project policies. Where inconsistencies arise, the lead Institute’s policy will be taken as binding.

Data Types and Sources. The data gathered and generated by the project will take two forms. “Science Driver” application data will be used as part of development and validation experiments and software artifacts (i.e. “code”) will persist as part of the research output. The application data that is needed to demonstrate reproducibility or that forms the basis of new domain science will be curated and stored at the participating institutions and the code will be managed via a reliable source-code control facility.

Data and metadata (‘data from now on’) produced by the research activities of the project will be shared and preserved. These data will be categorized in three levels:

- *Raw data* will be generated by experiments and further refined to be presented in publications. The repository of raw data will be backed up daily and preserved until at least after the end of project.
- *Wrangled data* will be a subset of raw data selected and produced via scientific analysis. Wrangled data will be shared via a dedicated and public git repository. Wrangled data will be preserved for at least two years after the project.
- *Curated data* will be the subset of wrangled data presented in one or more publications or other scientific products. Curated data will be shared via dedicated git repository, and will be preserved for five years after end of the project.

Both wrangled and curated data will be anonymized and no confidential or personal data will be publicly shared. At BNL, the data preserved in the context of this proposal will be further governed by National Laboratory policies pertaining to intellectual property and data management.

Standards for Data. The project will make use of standards for data and metadata whenever feasible. For example, GPS data is often given in NMEA format as standardized. Where existing formats are not available, open, or well-supported, we will default to using SI units and record the data in comma separated value (CSV) files in plain text and a relational database. Each file will start with a short annotation describing the fields. A detailed annotation will be stored in a README file that is stored along with the CSV file. This README file will also specify metadata that indicates how, where, and when this data was generated. To reduce data size and enable easier download, we may compress the CSV and README file together in a zip file or similar compressed archival format. We plan to use the CSV format as it is widely used and accessible on virtually all platforms; however, we will only use it as a fallback when other domain- or sensor-specific formats are unavailable.

Data will be stored in conventionally accepted digital formats for purposes of archiving, presenting, and publishing including tiff/tif, eps, ps, pdf, jpg/jpeg, png, gif, doc/docx, xls/xlsx, ppt/pptx, rtf, a variety of ASCII formats, mov, avi, wmv, flv, mp4, mp3 and others as found most efficient and suitable to each application. The project team will choose the appropriate format for the most-efficient conveyance, and maintenance of the quality and integrity of the data. For example, tiff and eps formats are best for publication-quality documentation for peer-reviewed journal submissions, while jpg and gif are suitable for inclusion in presentation-quality documentation for seminars and colloquia.

Data Sharing and Access. The core software we will develop will be made available via public repositories (e.g. github). Data will be made available online within 6 months of the results of the work being published in theses, dissertations, books, journals, conferences, etc. Data may be withheld for a reasonable amount of time if there are additional pending publications.

Publication, Dissemination, and Reporting. There will not be any restrictions placed on the use of project data, except in cases where distribution of raw data is deemed detrimental. The original authors and collectors of the data must be cited whenever the data is used and commercial gain from this data may be limited. Technical results published in conferences and journals will be archived by the relevant professional society library (e.g. ACM, IEEE) and also at the institutional digital libraries at BNL, UCSB, UNL, and Notre Dame

Data Preservation and Curation. Data, software, publications, and course material will be maintained and archived at UCSB and made available online. Archival data will be maintained for at least 5 years following the end of the grant period, although it is expected that the data will be maintained indefinitely by the team.

FAIR Principles. This plan achieves the FAIR principles as follows: *Findable.* Software and data will be searchable and findable by consumers via search interfaces (Github and Aristotle) and also indexed, linked, and described from the project web page. Both software and data products will be given descriptive metadata and assigned permanent DOIs as part of the publication process. *Accessible.* All software and data will be open source and directly accessible online without credentials. *Interoperable.* As noted above, the project makes use of standard formats for data to ensure interoperability with similar systems. *Reproducible.* The experimental work done by the project will be captured by completed archived scripts and environments that themselves point to named and archived code and data.

Software Licensing. The CSPOT software is made available under a permissive BSD License to (SPDX BSD-2-Clause) promote use and distribution of derivatives. Software will be stored on github to facilitate distribution and long term availability. In the event that some project works improves or contributes to outside software systems, those contributions will be governed by existing outside licenses.

BSSW Principles. Effective research in software systems requires that we strike a balance in development methods. On one hand, development must be undertaken with enough care to ensure that our products are of sufficient stability and quality to be used and have impact. On the other hand, conducting effective research requires that processes be lightweight enough that academic employees can focus on innovation. The PIs share a common philosophy of software development that involves the following principles:

- *Everything in the Repository.* The code repository contains the entire state of the project – source code, tests, examples, documentation – and nothing is considered to *exist* until it has been committed so that it can be shared and tested.
- *Integrated Testing.* Every commit (or pull request) to the repository results in a battery of automated tests to check the core functionality of the system. No commit is accepted that fails the tests. Every expansion of functionality must include relevant tests.
- *End-to-End Deployment Testing* The entire deployment process of the software – installing dependencies, building code, downloading data, invoking the system – is captured as a script that forms part of the public documentation consumed by end users. Then, this documentation is executed daily as an end-to-end test that the documentation itself is a correct and complete path to a working output.
- *Design Changes Before Coding.* We teach our students and staff to write design changes documents prior to touching code. These need not be elaborate, and are often just a few paragraphs describing a perceived problem, an overview of one or more solutions, and some notes on the scope of the changes needed. We find that these documents encourage other project members to chime in and point out potential problems and alternate solutions, thus avoiding technical dead ends and conflicts.

Appendix 5 – Promoting Inclusive and Equitable Research (PIER) Plan

Our goal with this PIER plan is to promote diversity, equity, inclusivity (DEI), and accessibility of the scientific enterprise – *through* XGFABRIC, and to use our research program to broaden participation in science, technology, engineering, and math (STEM) by students from underrepresented minority (URM) groups. To enable this, we propose to activities that focus on (i) recruitment of students from diverse backgrounds, (ii) creating and sustaining an inclusive, professional, and supportive collaborative research environment across our institutions, and (iii) training and mentoring for the development of a diverse workforce. Our team is uniquely positioned to pursue these activities given our successful track record with outreach, mentoring, training, and increasing representation of students and researchers from diverse groups – both at the institutional level and individually. Specifically, we have mentored and advised female and URM students who have gone on to do exceptionally well in STEM, with successful careers and exciting lives in both industry and academia. We also have pursued novel outreach activities for K-12 students, first-generation college students, and students from financially disadvantaged households. Our experience includes training (e.g. after school STEM academies, scientific classification of animals), mentoring (e.g. college visits and advising of students at all levels), socializing (e.g. field trips to see scientific instruments and multi-scale IoT and atmospheric science deployments), and activities that build confidence in their technical and scientific abilities (e.g. “hack-a-thons” for digital agriculture applications). Recently, our team has worked with students from rural communities as part of collaborations with community colleges and universities with teaching mandates, to provide training in computing and science basics, as well as hands-on experiences with multi-scale, distributed IoT systems. As we describe in Section II, our PIER plan for this project leverages and builds upon these efforts but is tailored to and tightly integrated into our XGFABRIC research plan. This will enable us to enhance our innovation while overcoming obstacles to achieving equity in computing.

Section I – Leveraging Institutional-Level Programs

Our PIER plan will leverage an extensive foundation of institutional-level programs and successes at each of the PI’s institutions. All three institutions actively integrate DEI as part of their overall operations and work closely with MSIs via a mix of projects and mentoring relationships. For example, UCSB was the first member of the Association of American Universities (AAU) – a prestigious group of tier one research universities (R-1) – to earn the Hispanic Serving Institution (HSI) designation – and thus has a long track record of successful recruitment, training, and mentoring activities.

An inclusive environment and diverse workforce are also essential for Brookhaven Lab to achieve its mission on behalf of the DOE’s Office of Science. BNL promotes equity to ensure that historically underrepresented or underserved people in STEM have opportunities to grow and contribute in their chosen fields by applying consistency in employment practices and promoting collaboration and fairness. BNL’s DE) Office supports the mission by facilitating the development pipeline of qualified, diverse candidates; safeguarding equitable treatment for all employees; and supporting an environment that encourages respect for individual differences through inclusivity. In addition, BNL’s Talent Acquisition team, DEI Office, and Office of Educational Programs (OEP) collaborate to establish pathways for undergraduate and graduate students through workforce and peer development programs while working closely to ensure next-generation STEM leaders will be diverse, equitable, and inclusive.

Similarly, Notre Dame promotes DEI through academic inquiry, and numerous programs and opportunities that reflect their foundational belief in the dignity of all men and women. Their programs and strategies for building DEI flows from the University’s commitment to intentionally creating an environment of mutual respect, hospitality and warmth in which none are stranger all may flourish. As with UCSB and BNL, Notre Dame as an institution has extensive experience partnering with MSIs and non-R-1 institutions to build and sustain as strong STEM/STEAM pipeline and workforce.

As a founding member of the BRAID initiative (Building, Recruiting, And Inclusion for Diversity) , co-led by AnitaB.org and Harvey Mudd College, the University of Nebraska-Lincoln’s School of Comput-

ing has a long track record of improving participation of underrepresented minorities (URM) and female students. As a part of the departmental broadening participation in computing (BPC) plan, UNL projects and activities are organized under three primary sustainable, coordinated categories: (1) Recruitment (RC) activities to broaden the pipeline of undergraduate students, (2) Retention (RT) efforts to increase degrees conferred to diverse students, and (3) Community-building (CB) activities to facilitate peer-support for students on and off campus. We will participate in existing research, training, and outreach activities across our institutions and engage, mentor, and train participants regardless of their gender, race, disability status, or sexual orientation. Specifically, we propose to leverage and build upon this institutional level DEI foundation to implement informed, effective, and high impact actions as part of our PIER plan for XGFABRIC.

Section II - Planned Actions

We propose PIER planned actions that target DEI-aware recruiting, retention, and training. This will involve pursuing activities individually and collaboratively to enhance and sustain the STEM pipeline for URM and female students and create best practices for an inclusive environment for STEM collaborations.

Recruitment and Retention

The PIs will each pursue STEM pipeline building activities by leveraging the programs at their institutions and tailor each experience to incorporate XGFABRIC advances and applications. At UCSB, the PIs will mentor teams of students from the Early Research Scholars Program (ERSP) [10] to engage and retain female and URM undergraduates in computer science. ERSP provides a team-oriented, research apprenticeship to second-year female and URM undergraduates. Our team will mentor 1+ teams from this program and craft their projects according to their skill and experience level. The ERSP students will work closely with our graduate students for the most rewarding experience possible. Specifically, we will use XGFABRIC as part of each project to give students hands-on experiences with new programming and distributed systems with which they can develop, debug, and reason about the next-generation of adaptive and resilient applications and workloads. We believe that doing so will help build student problem-solving skills and confidence with technology, and give students a broader view of computing and its potential for impacting society.

At BNL, we will extend our Computational Science Initiative (CSI), which actively integrates diversity and inclusivity as part of its research effort and overall operation. CSI has long maintained connections to MSIs, historically black colleges and universities (HBCUs), and HSIs via a mix of active projects and mentor relationships. We plan to host graduate and undergraduate interns annually through various local and national programs including the Sustainable Research Pathways for High-Performance Computing program, a partnership between DOE and Sustainable Horizons Institute, to provide a pipeline to the national labs for scientists and engineers from underrepresented communities. Students will be taught/mentored on XGFABRIC advances and subprojects including workflow management and system resiliency as part of their internships. Moreover, to encourage career growth, we will develop professional development activities through CSI, specifically geared toward postdoctoral researchers and early-career scientists (supported by the XGFABRIC project and elsewhere) that provides foundational skills and tools in scientific and distributed programming for multi-scale systems while they are working and being mentored.

At Notre Dame, we will help develop and participate in programs that better prepare undergraduates for computing as part of a new summer program (called ELITE) for programming novices. The program focuses solely on developing programming experience and technical skills through a series of leveled exercises, combined with regular networking and career development opportunities. This project will interact with ELITE in two ways: (1) supported graduate students will serve as programming and career mentors for students in the ELITE program; (2) summer undergraduate students funded by this project will be recruited *from* the ELITE program to participate as summer research assistants funded through this program, following the sophomore year; and (3) Those who have gained experience in the project techniques through Notre Dame will be encouraged to apply for DOE SULI internships. Through these activities, we will develop a pipeline from initial programming instruction to research participation.

UNL's recruitment and retention activities will include recruiting students through UNL SoC Grad Info Day, Hour of Code, Sunday with a Scientist. UNL SoC Grad Info Day is a two-decade-long recruitment event that invite prospective Midwestern students to a two-day recruitment activities. During these activities, we will present and showcase the XGFABRIC results to recruit students to the project. Similarly, these activities include demos and tours of labs that are led by graduate students, creating a sense of belonging and ownership that is essential for retention. Hour of Code is an annual event organized by the Lincoln Public School system, in collaboration with the UNL SoC, that houses over 600 elementary and middle school students for a day of activities related to computing. We will continue to reserve a booth at this event to showcase IoT and wireless communication basics through hands-on activities. Sunday with a Scientist is a monthly event hosted by the Nebraska State Museum that invites a scientist for a day-long public interaction with Lincolniters. We will present XGFABRIC results to disseminate project results to the public.

Regular PI meetings will also provide students and staff the opportunity to share and showcase their accomplishments, giving the other institutions/PIs and industry participants the opportunity to recruit them (for graduate work, post-doctoral research, teaching/training in STEM positions, industry positions, and research internships). Finally, we will assess our impact via interviews before and after, which query the understanding of the role of computing and science in their lives, as part of their educational path, and on society. Our metrics for success will be the relative improvement or broadening that we observe in this understanding, as well as student retention and placement.

Enhancing the Research and Training Environment

This part of our PIER will support the environment that surrounds collaborative research. To enable this, we will construct a XGFABRIC mission statement and code of conduct together and use it to construct a training plan that we can use to teach students and staff the importance of DEI in the pursuit of research and workforce development. We will leverage the DEI training experts from each of our institutions to provide instruction. To tailor these activities to XGFABRIC, we will use our target applications (climate change, energy independence, agriculture) as use cases to discuss the importance of considering environmental justice, end-user feedback, societal impact, and unintended consequences as we pursue our technical and scientific advances. We will also encourage and support graduate student participation in professional development, including annual diversity and scientific conferences, e.g., the Richard Tapia Celebration of Diversity in Computing and Supercomputing.

To ensure that PIs and staff succeed at mentoring, we will develop training that leverages XGFABRIC. Often, faculty and graduate students participate in mentoring and outreach activities with little experience or understanding of "how to mentor". With this activity, we will work with the UCSB Center for Science and Engineering Partnerships and the Office of Education Partnerships, and the BNL, UNL, and Notre Dame mentorship training professionals to invite professionals with to teach our mentors modern mentoring and DEI practices, and to provide each with suggestions and feedback on their mentoring efforts related to XGFABRIC. Our goal is to ensure that every mentor is impacting students to the best of their ability. We will assess our impact via participant interviews before and after, which query the understanding of the role of a mentor and best practices for mentoring and DEI in a computer science and computer engineering contexts. We also plan to interview student teams (mentees) before and after (one quarter later) their mentors participate in this training. Our metric for success will be the relative improvement in knowledge and understanding about techniques we can use to be good mentors, and relative improvement in the views of mentees in the program. We plan to open this activity up to our departments and invite others to share and participate in this workshop event as we gain experience with it, to maximize our impact. Finally, we will perform biannual surveys to assess and enhance the impact of our PIER activities. We will perform surveys at each institution and then discuss the results together to determine how to implement any changes. We plan to disseminate our results and findings via PI presentations and publication in appropriate seminar series and other venues, so that others can benefit from what we learn.

Appendix 6 – Other Attachment

This page intentionally left blank.